

Deep Learning Frameworks

Non-Linear Boundaries, Activation Functions, Backprop, Autograd,
Dataset, DataLoader and Regularization

By Sakharam

<https://tinyurl.com/dlframeworks>

<https://github.com/sakharamg/DeepLearningFrameworks>

Recap

- Computation Graph
- Tensor and Operations
- Toy Dataset – Initializing Tensors (Random, Zeroes, Ones)
- Neural Network – nn.Sequential
- Training Loop - Skeleton

<https://tinyurl.com/dlframeworks>

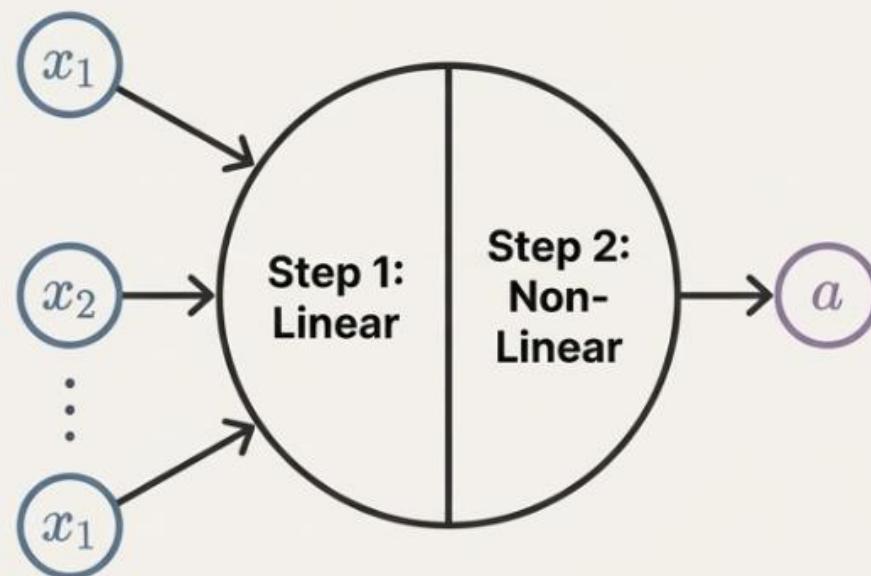
<https://github.com/sakharamg/DeepLearningFrameworks>

Neuron: Building Block of Neural Network

Part 1: The Linear Step

First, the neuron computes a weighted sum of its inputs, and adds a bias. This is a familiar linear transformation.

$$z = Wx + b$$

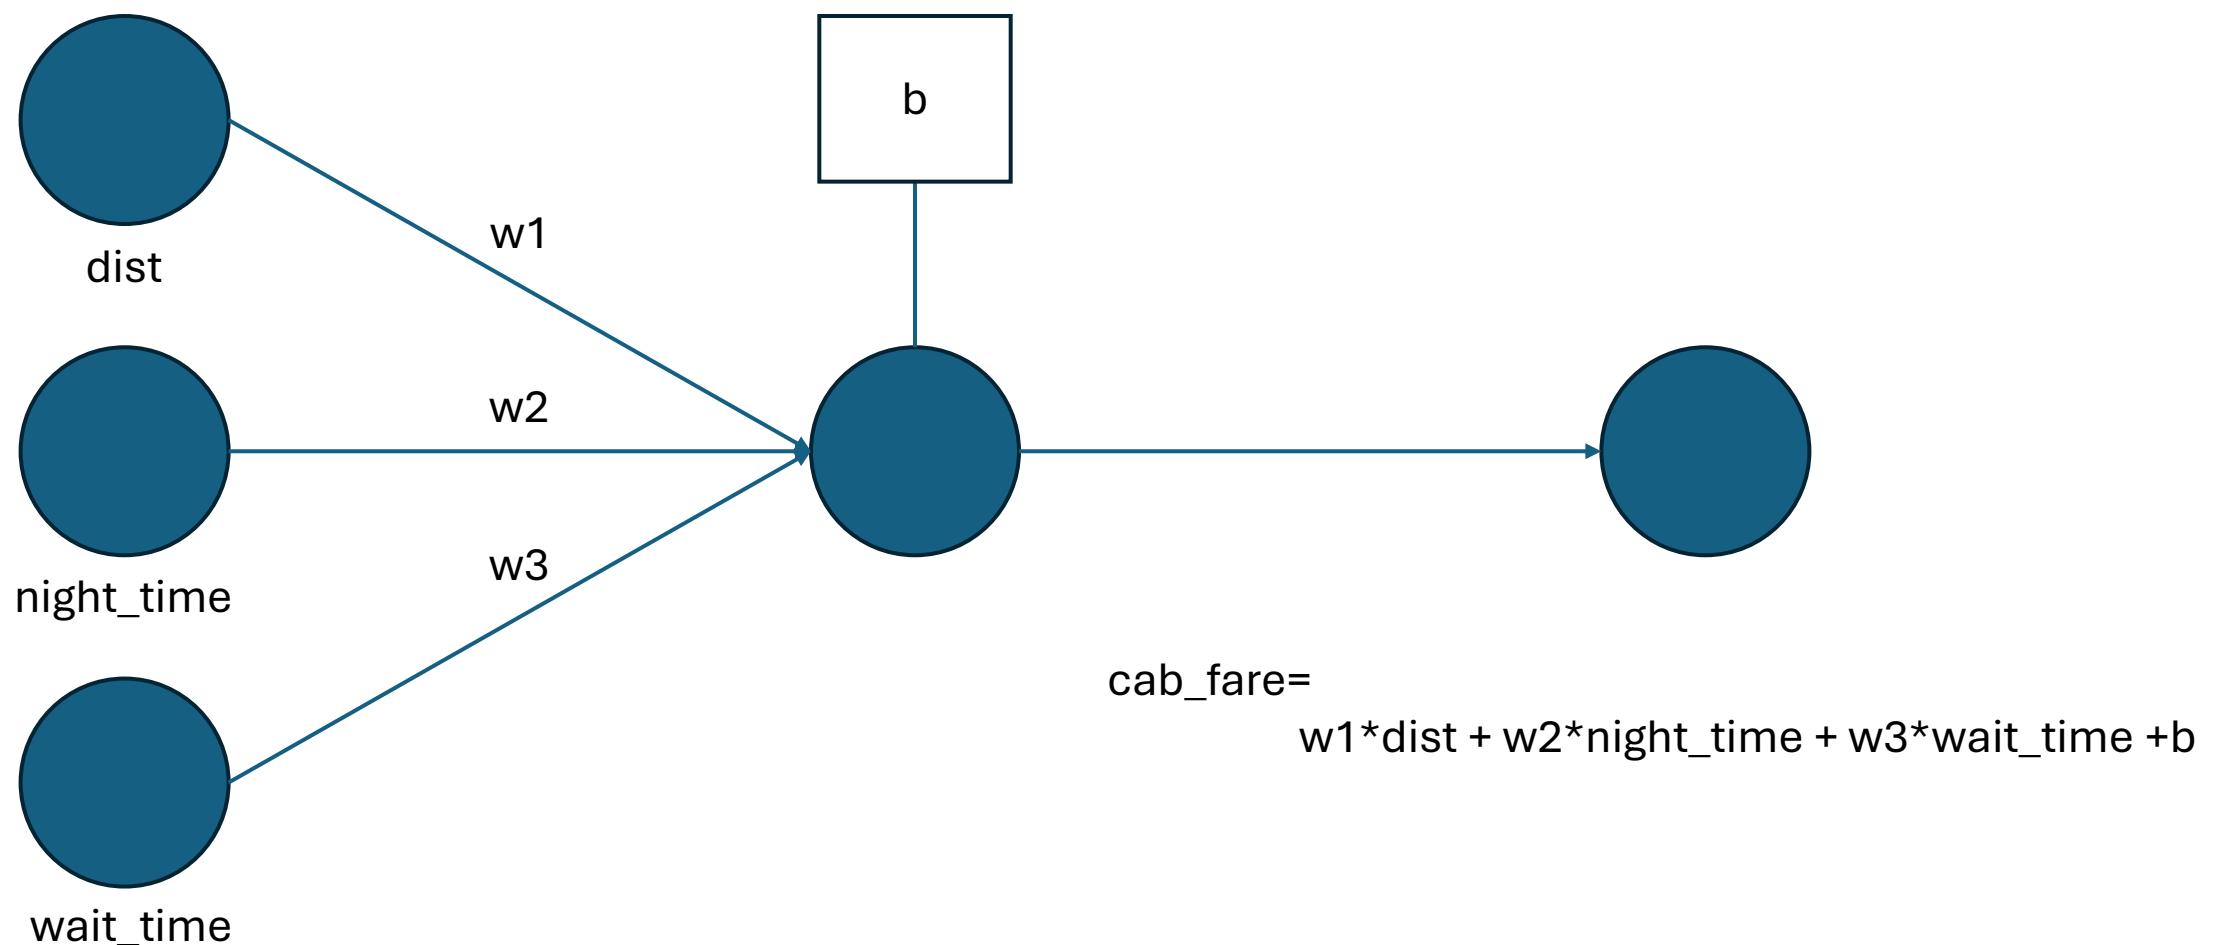


Part 2: The Non-Linear Step

Second, it passes this result through a non-linear "activation function".

$$a = f(z)$$

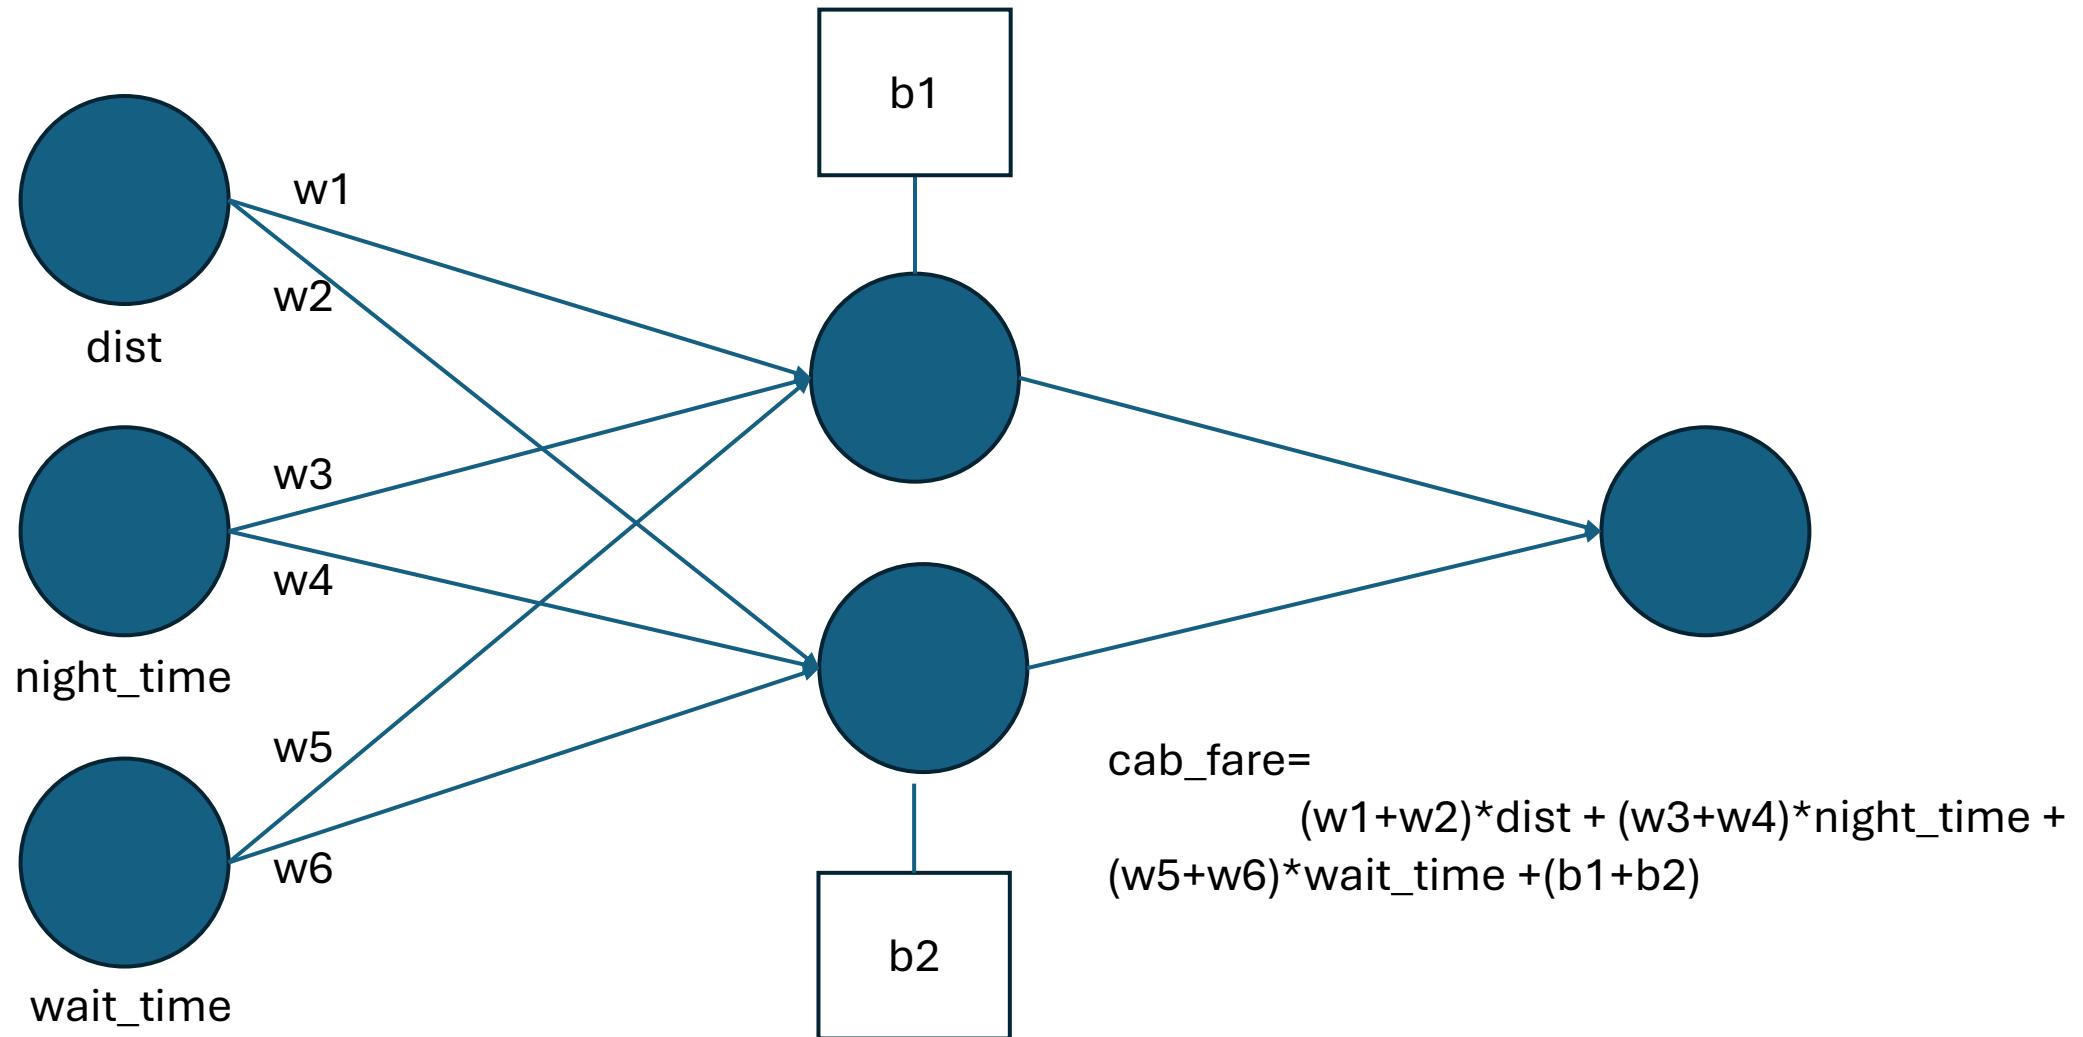
Neural Network with Linear Neurons



<https://tinyurl.com/dlframeworks>

<https://github.com/sakharamg/DeepLearningFrameworks>

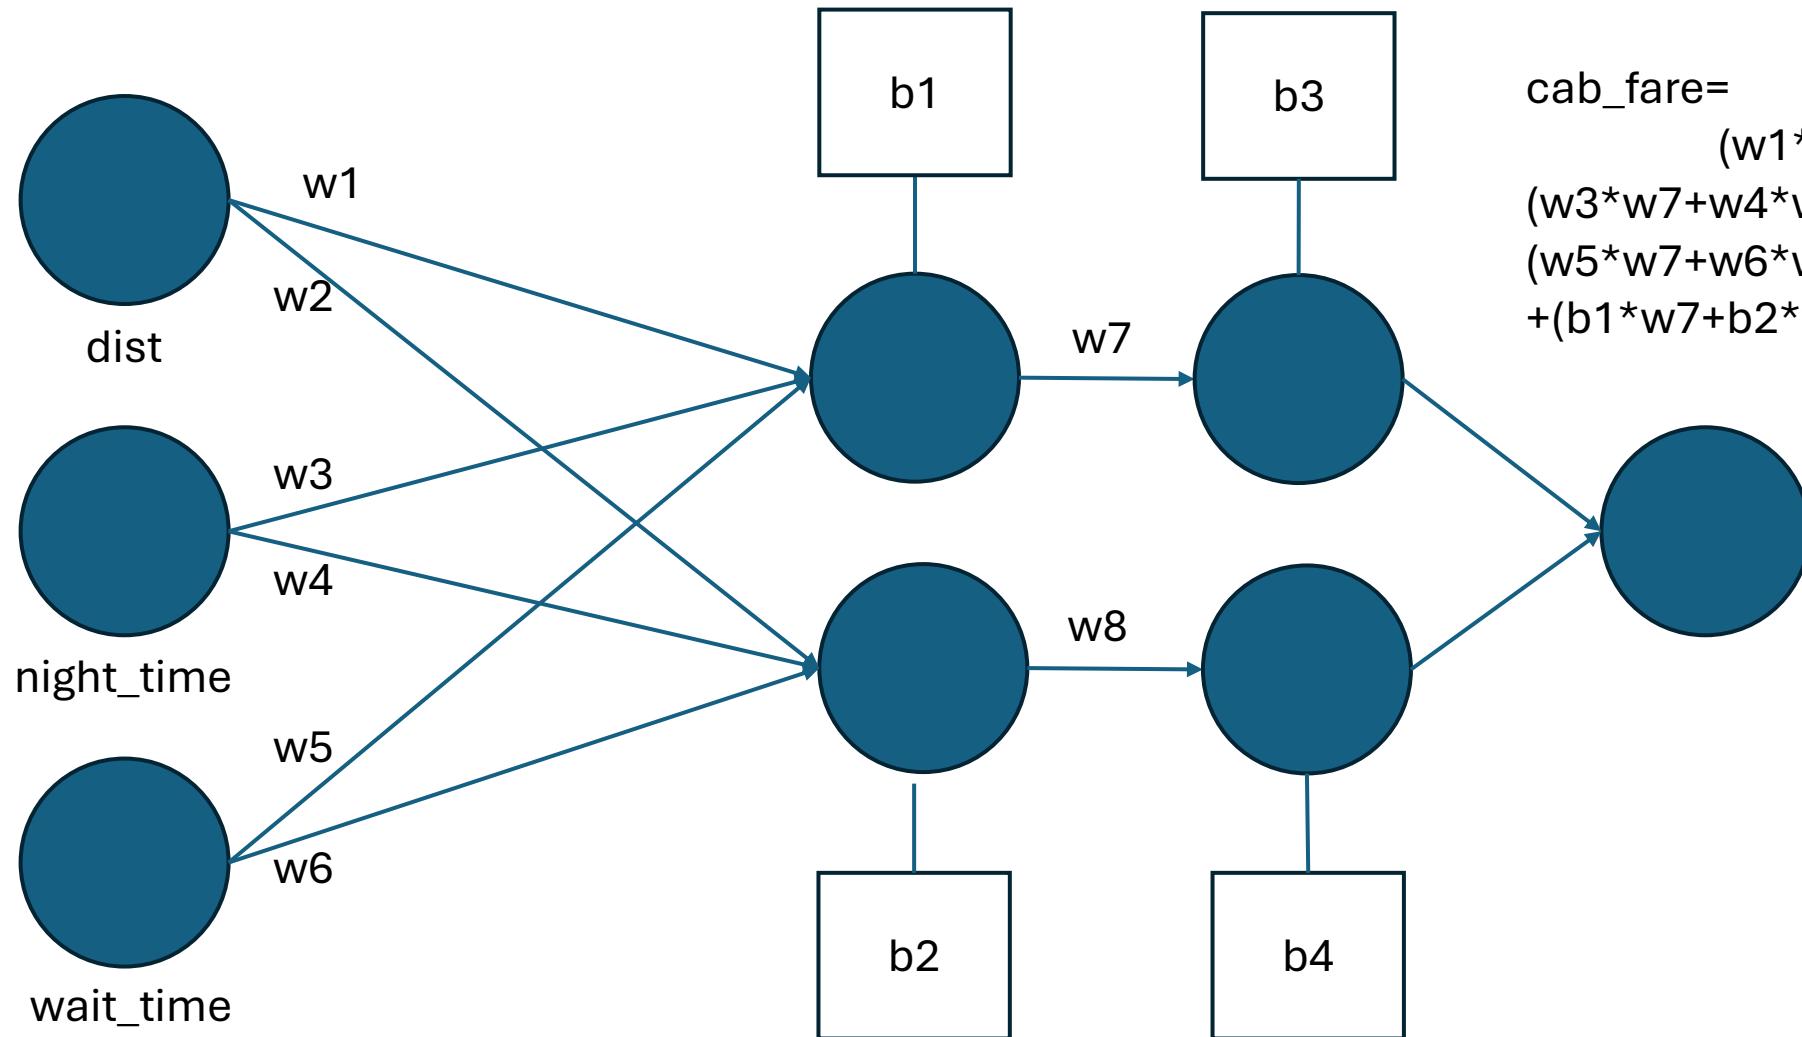
Neural Network with Linear Neurons



<https://tinyurl.com/dlframeworks>

<https://github.com/sakharamg/DeepLearningFrameworks>

Neural Network with Linear Neurons

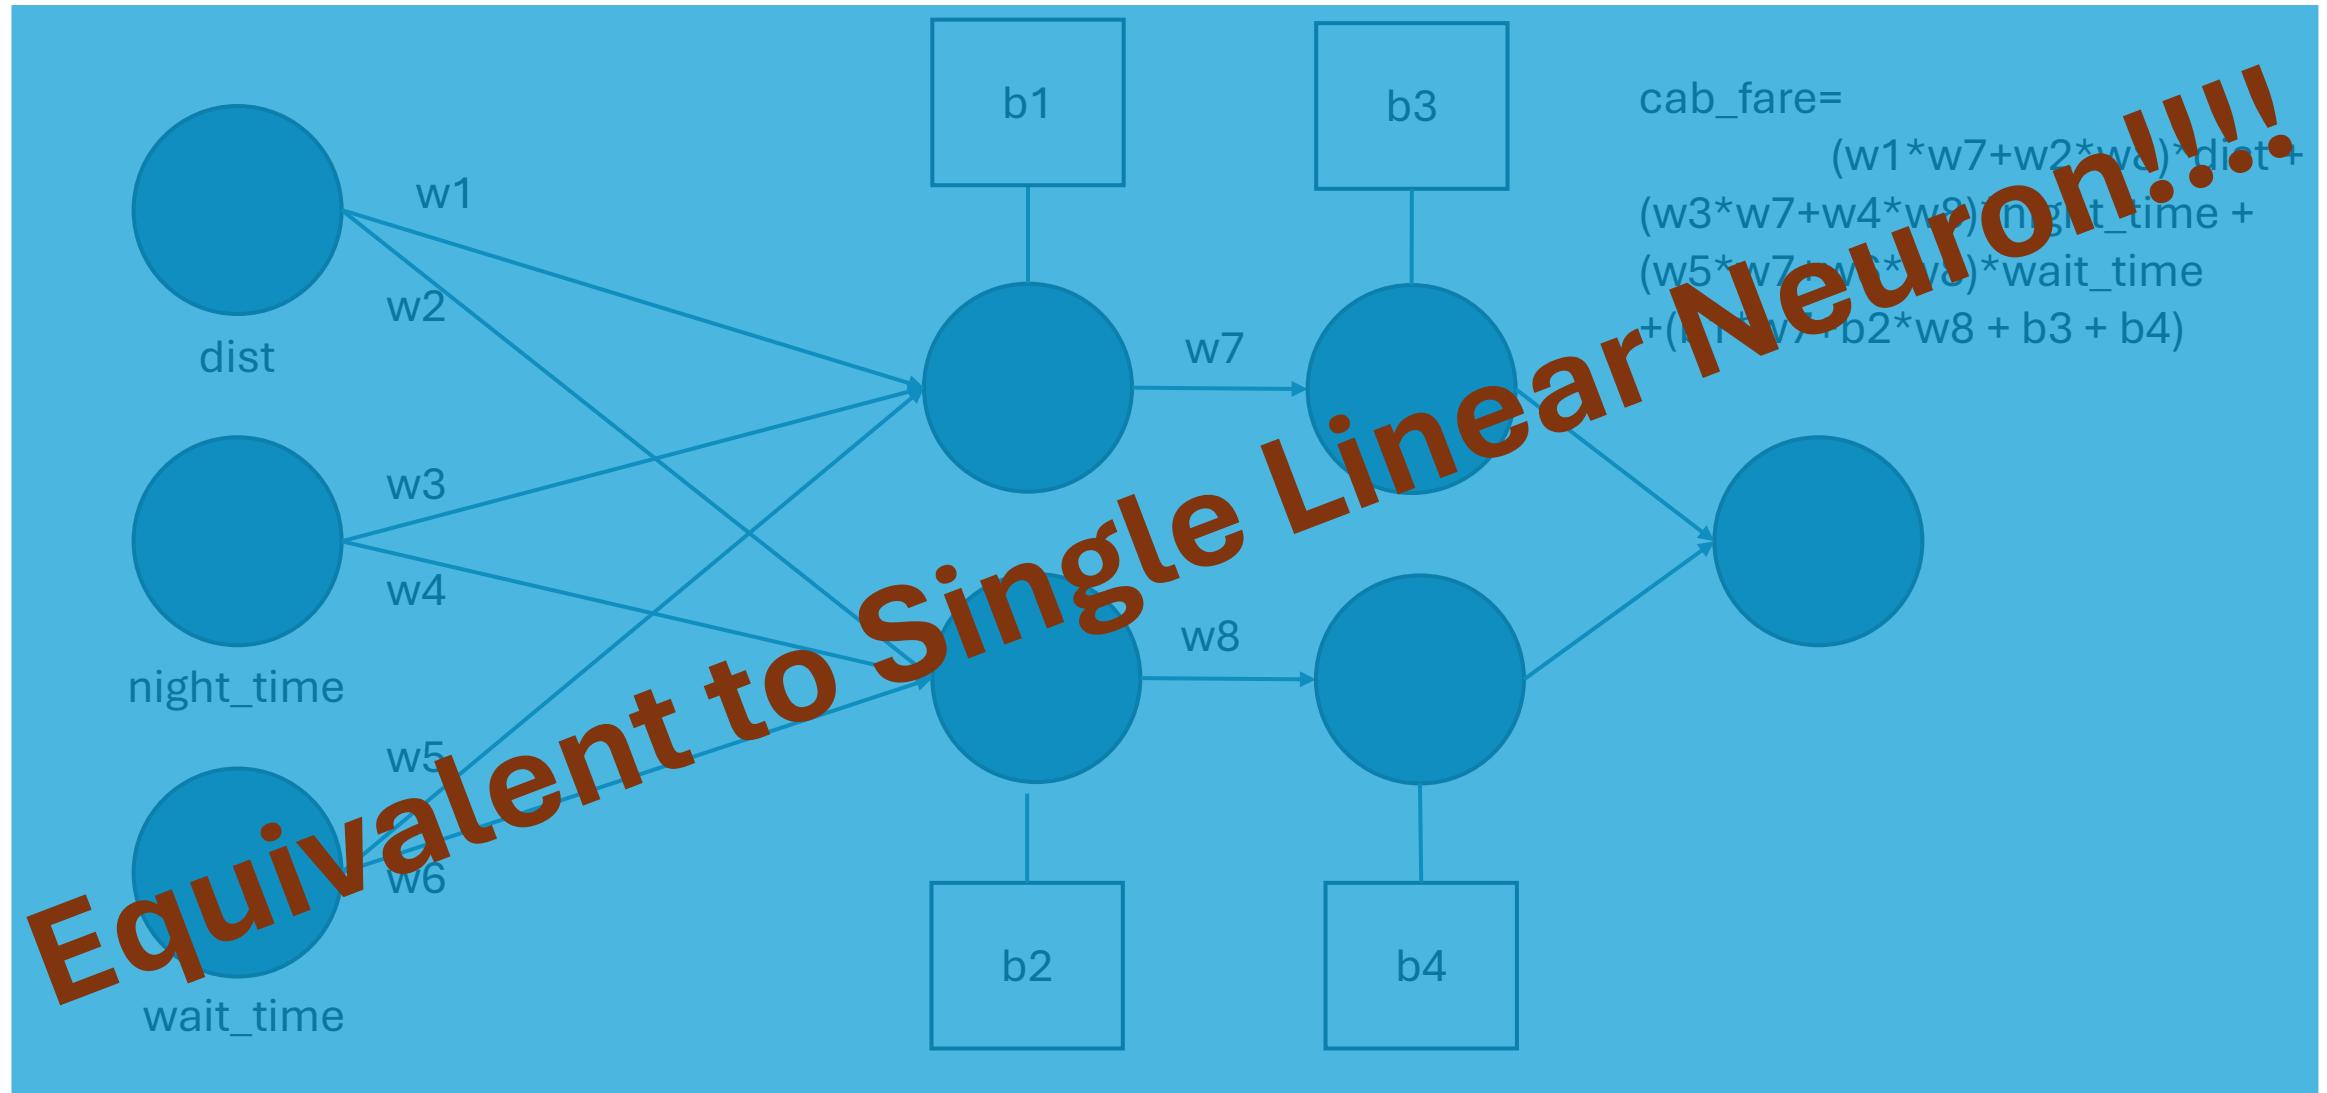


$$\text{cab_fare} = (w1 \cdot w7 + w2 \cdot w8) \cdot \text{dist} + (w3 \cdot w7 + w4 \cdot w8) \cdot \text{night_time} + (w5 \cdot w7 + w6 \cdot w8) \cdot \text{wait_time} + (b1 \cdot w7 + b2 \cdot w8 + b3 + b4)$$

<https://tinyurl.com/dlframeworks>

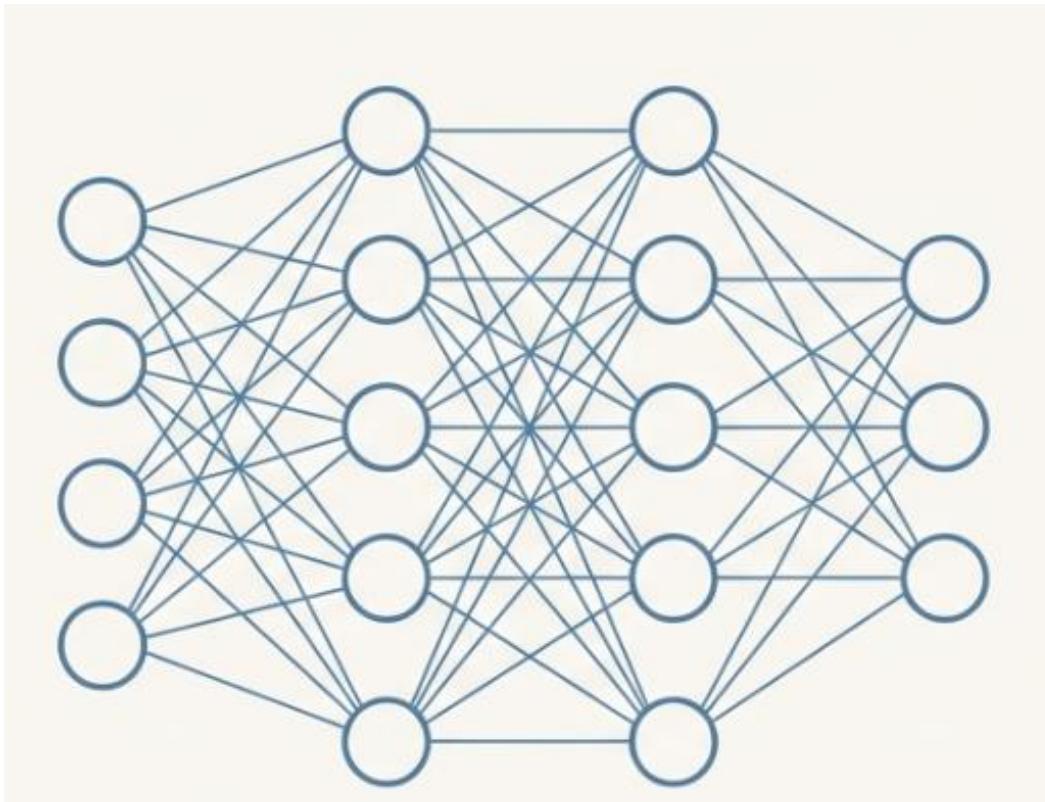
<https://github.com/sakharamg/DeepLearningFrameworks>

Neural Network with Linear Neurons



Neural Network with Linear Neurons

Still Linear



<https://tinyurl.com/dlframeworks>

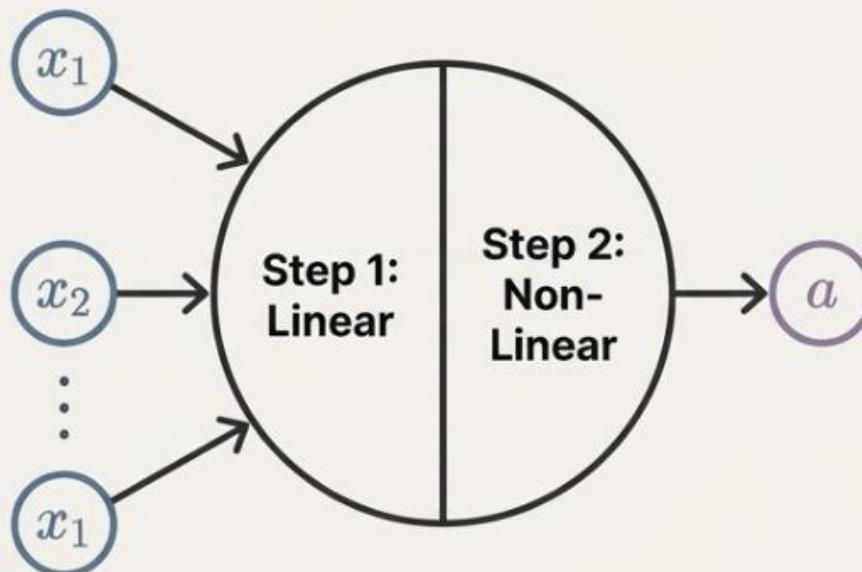
<https://github.com/sakharamg/DeepLearningFrameworks>

Activation Functions – Non Linearity

Part 1: The Linear Step

First, the neuron computes a weighted sum of its inputs, and adds a bias. This is a familiar linear transformation.

$$z = Wx + b$$

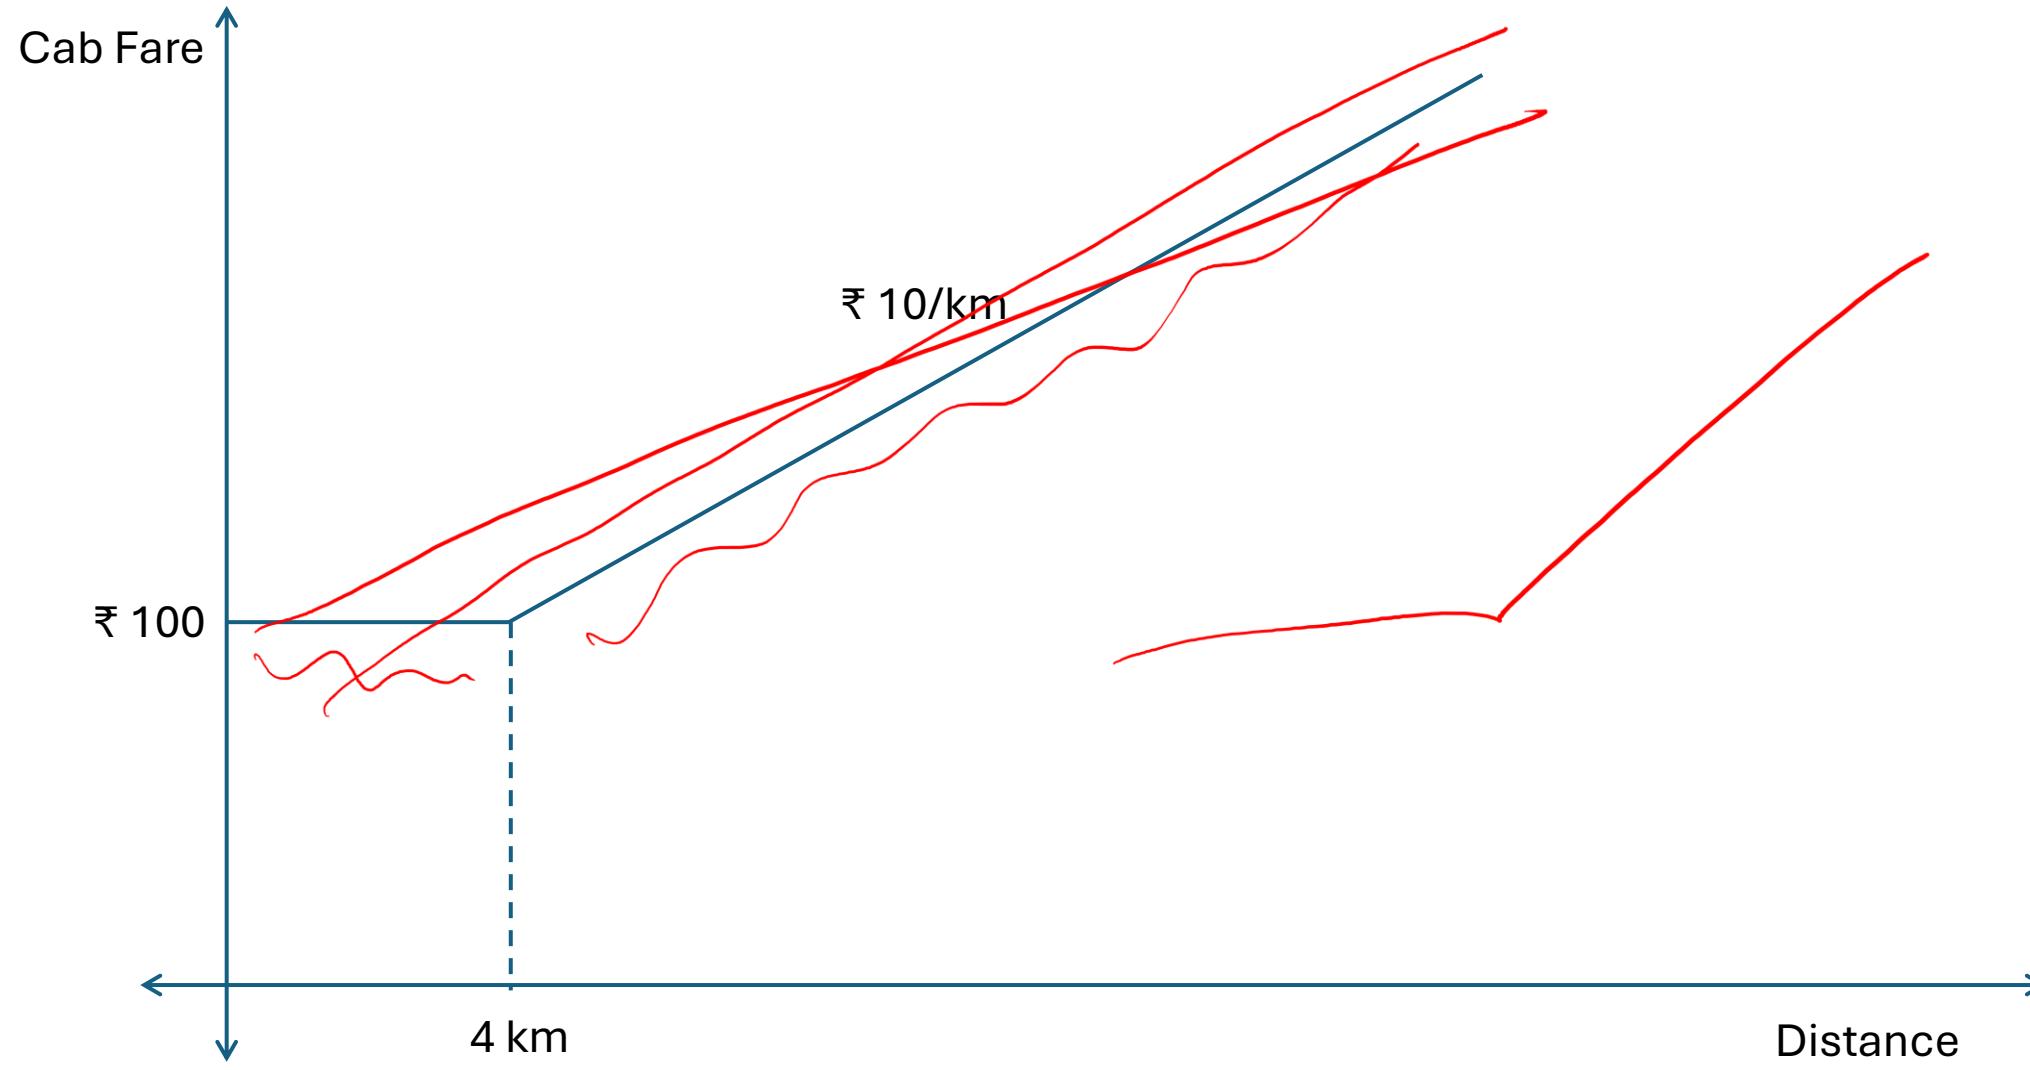


Part 2: The Non-Linear Step

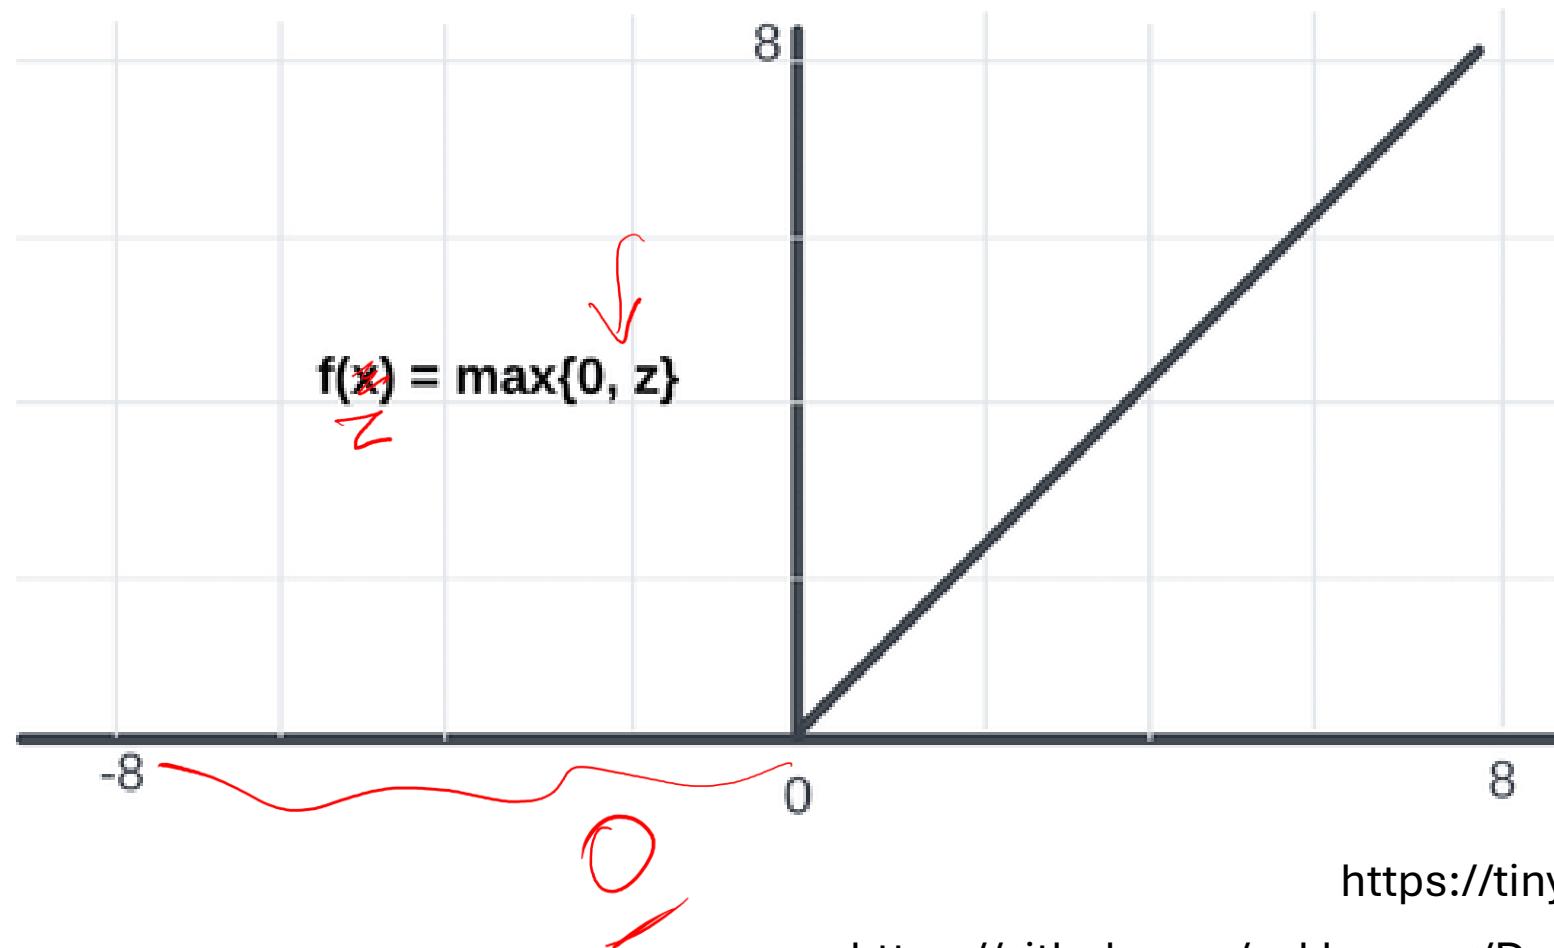
Second, it passes this result through a non-linear "activation function".

$$a = f(z)$$

Towards Non Linearity – Handling Bends

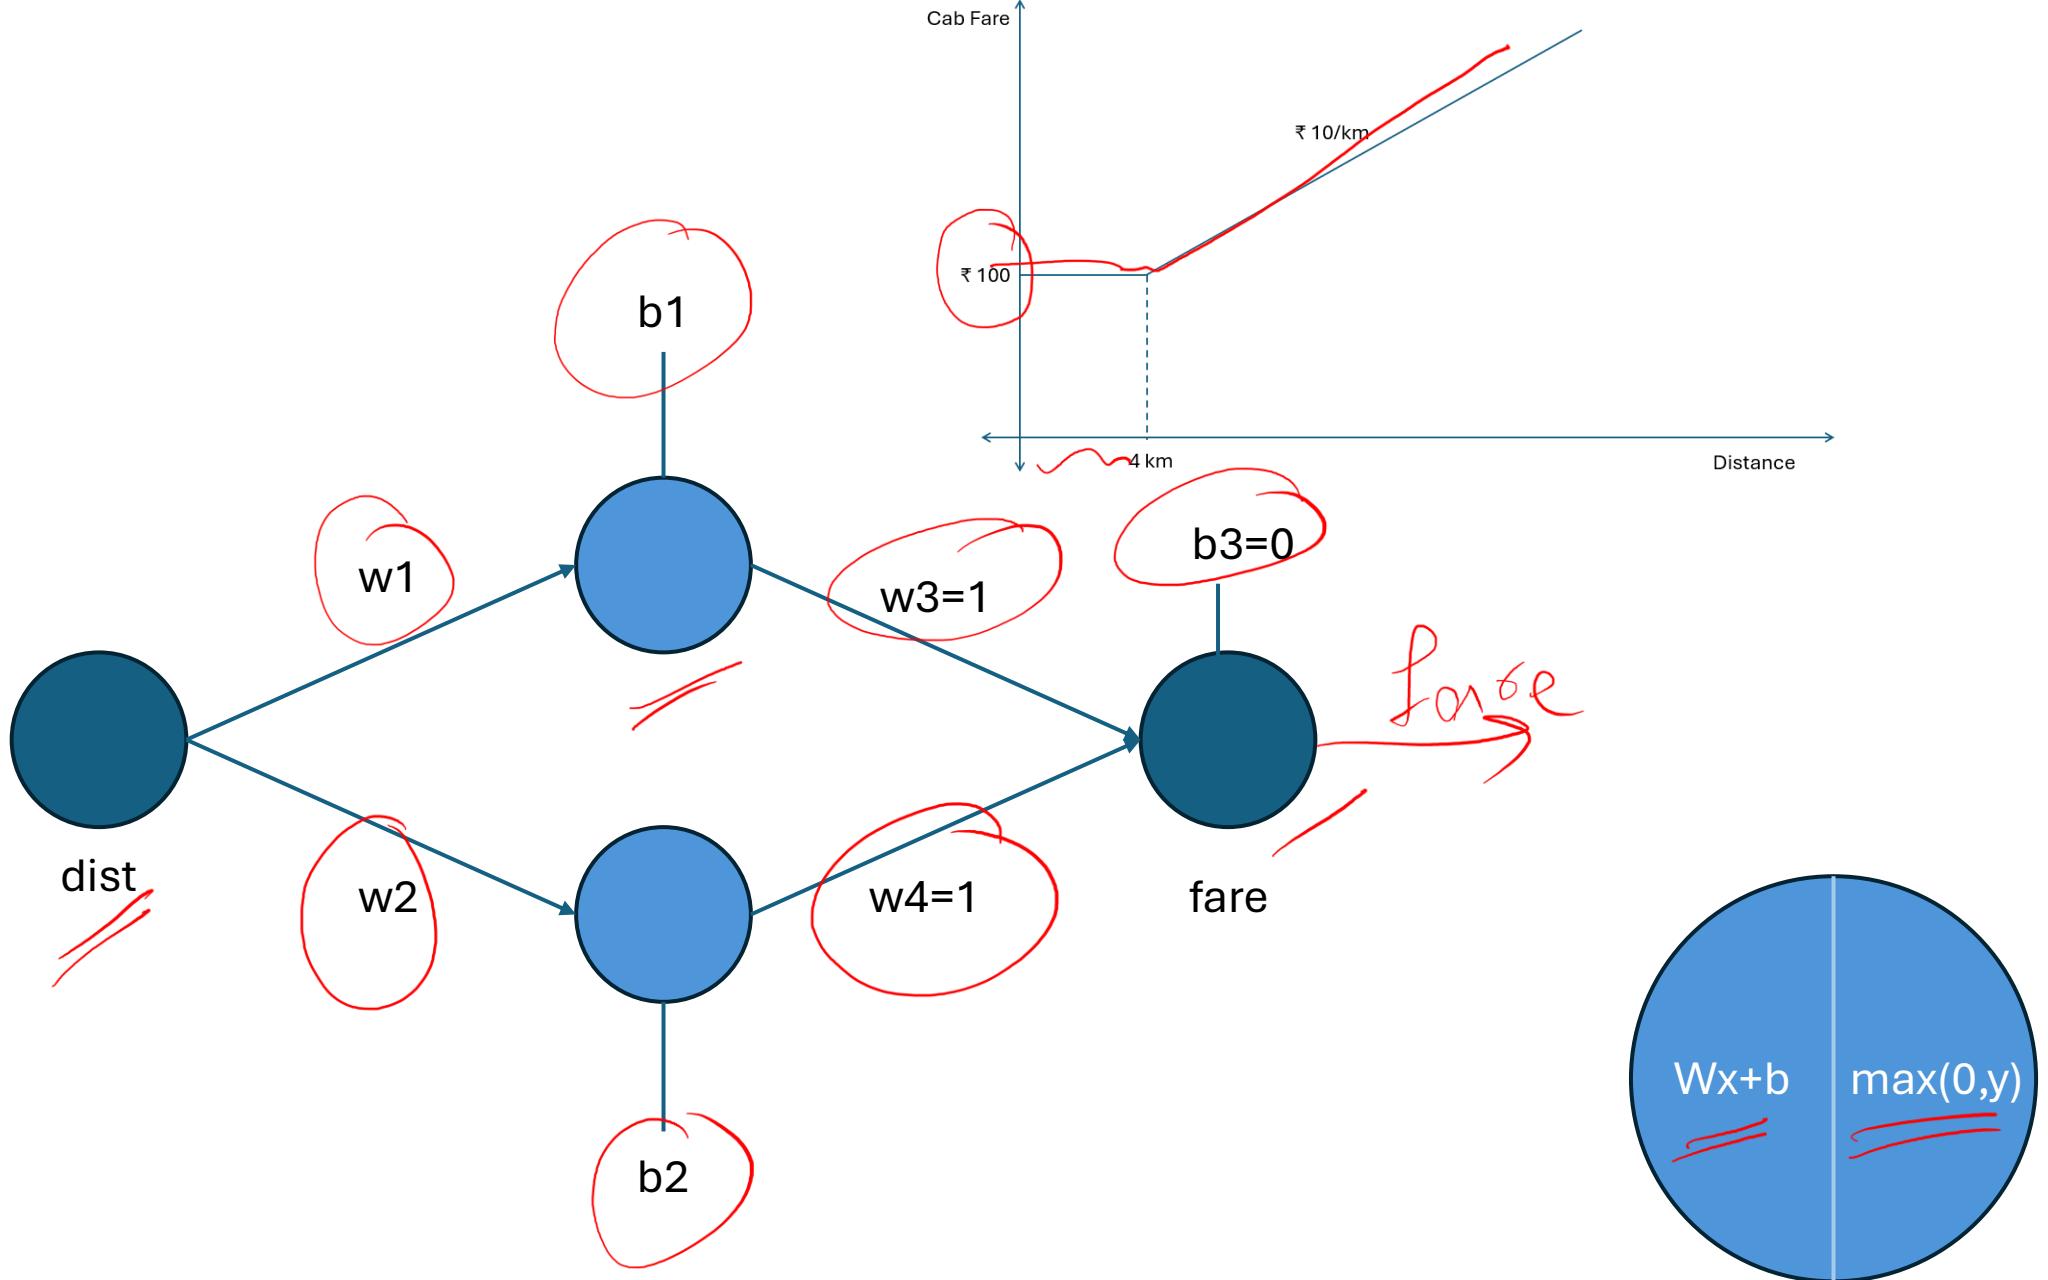


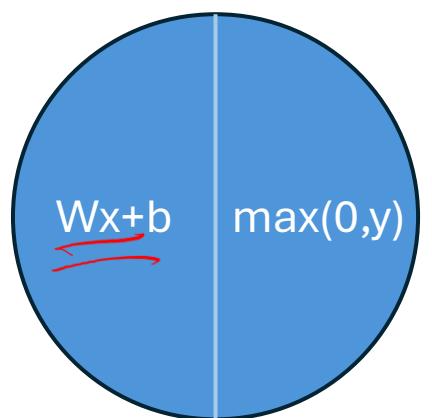
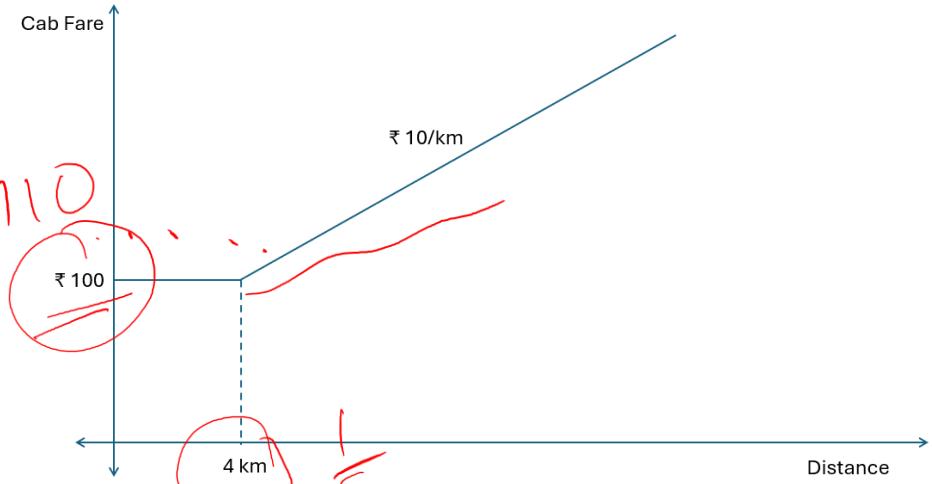
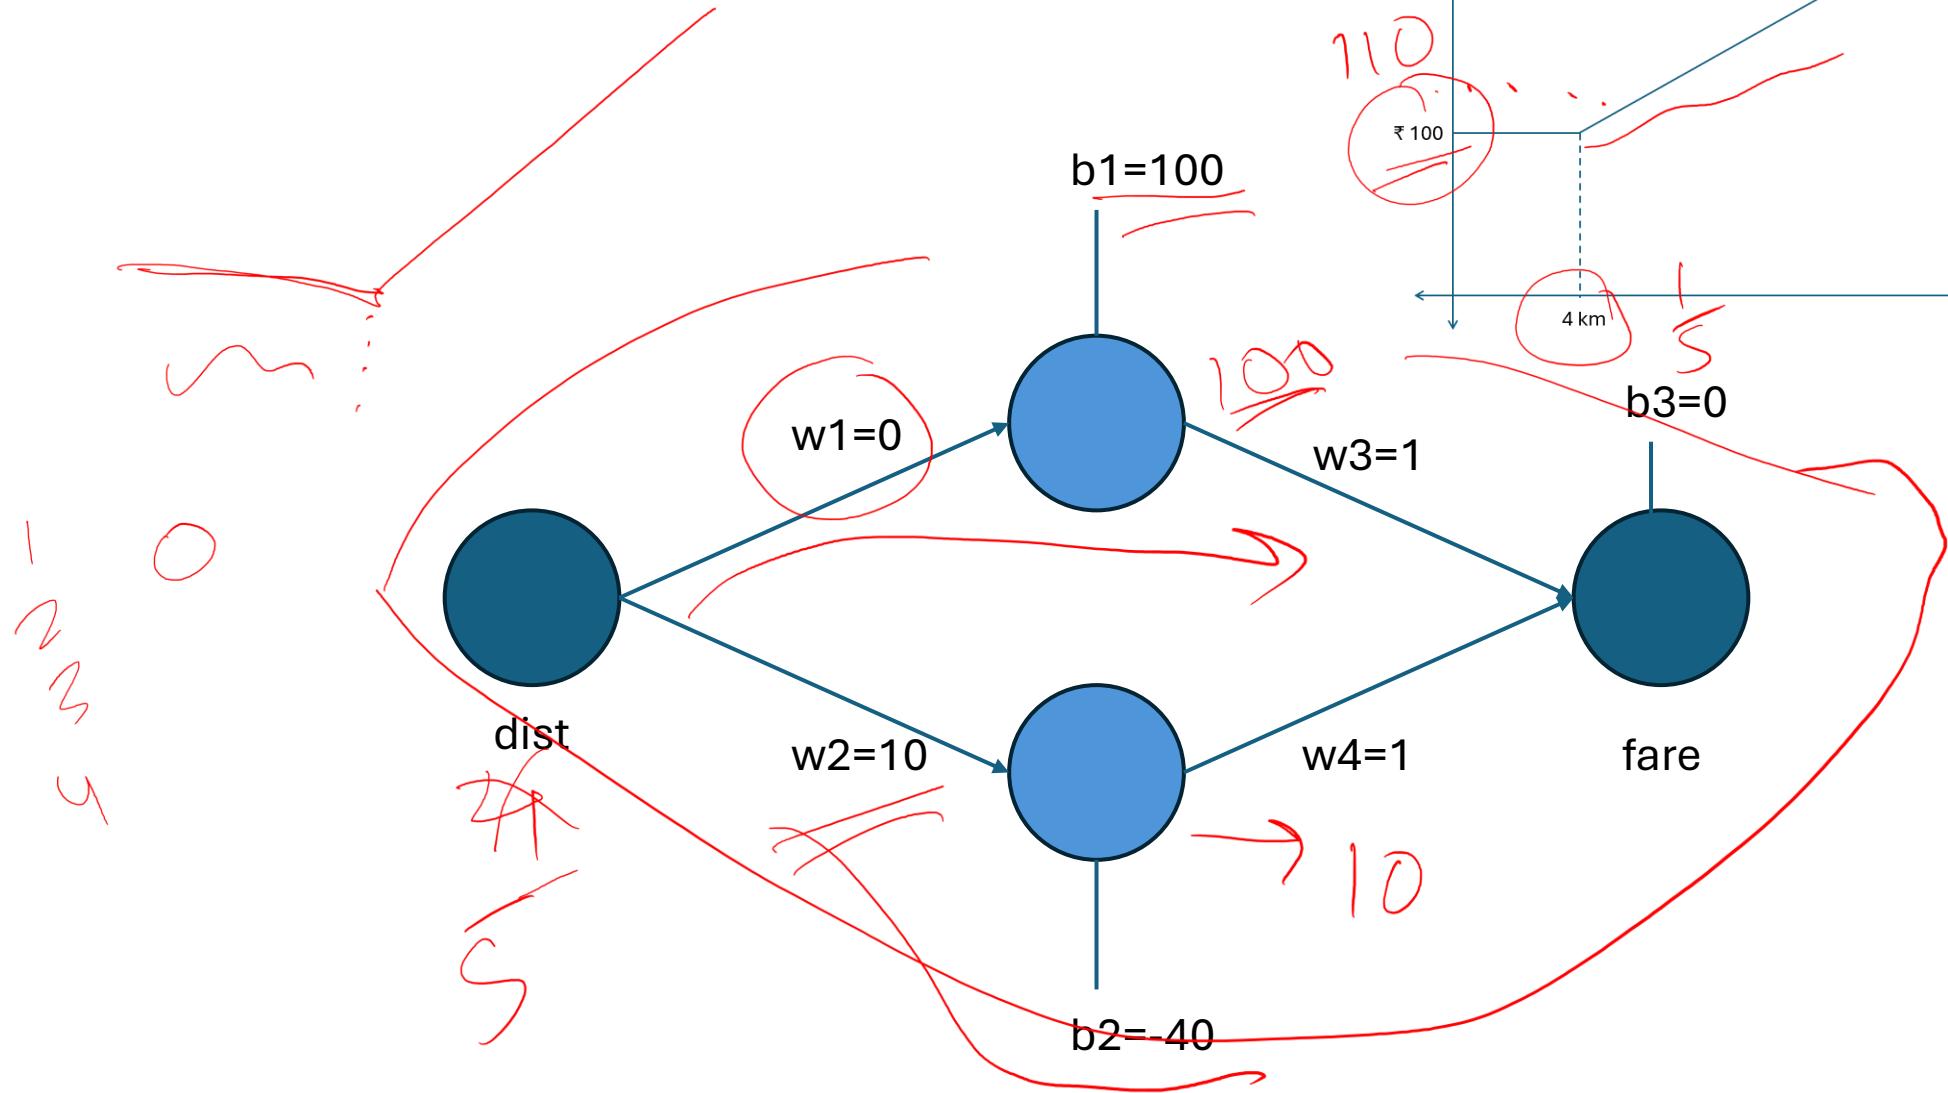
ReLU



<https://tinyurl.com/dlframeworks>

<https://github.com/sakharamg/DeepLearningFrameworks>





Classification

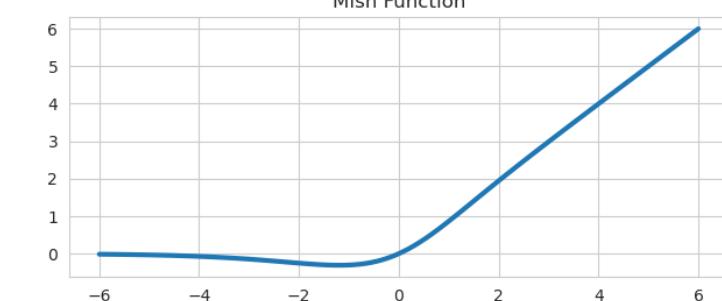
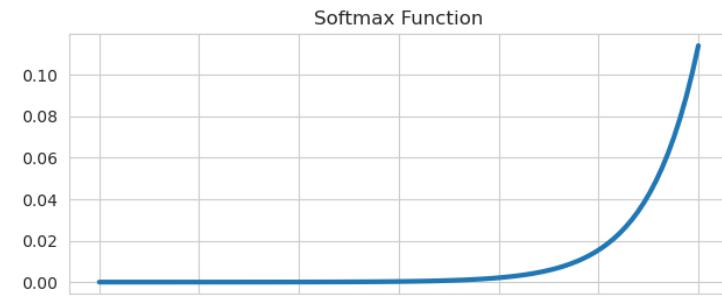
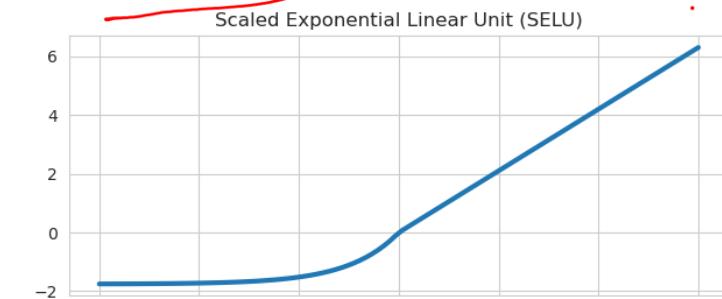
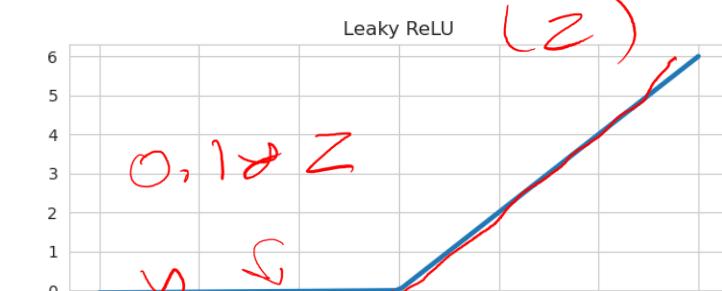
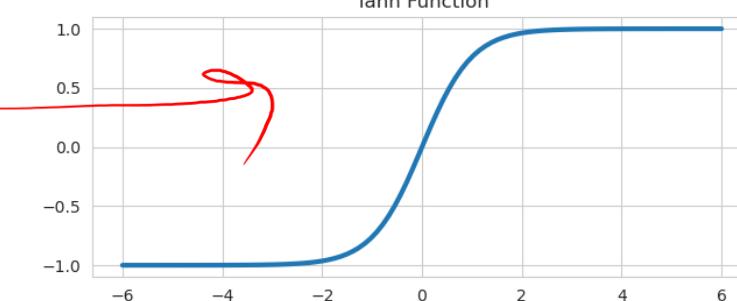
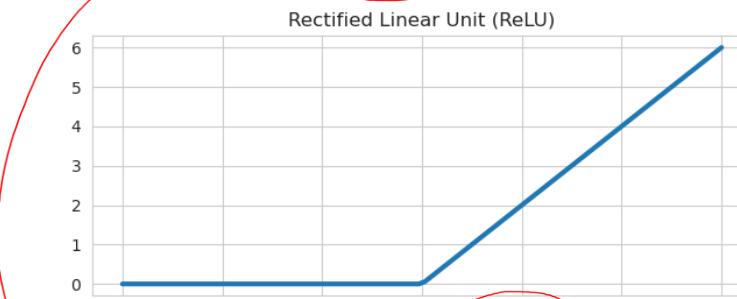
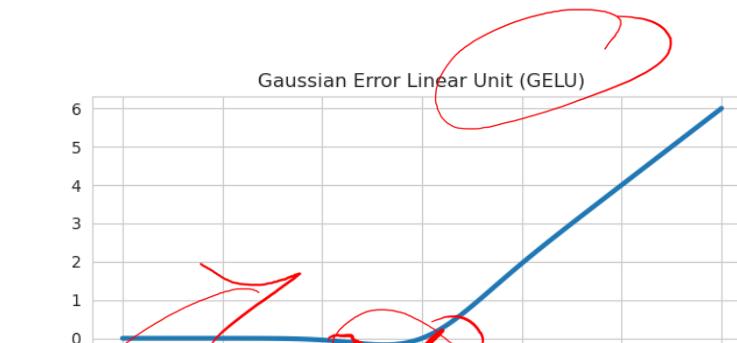
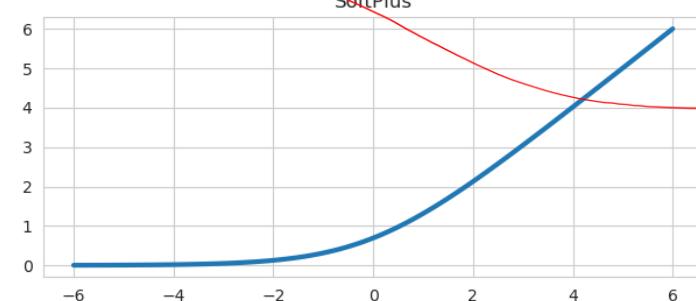
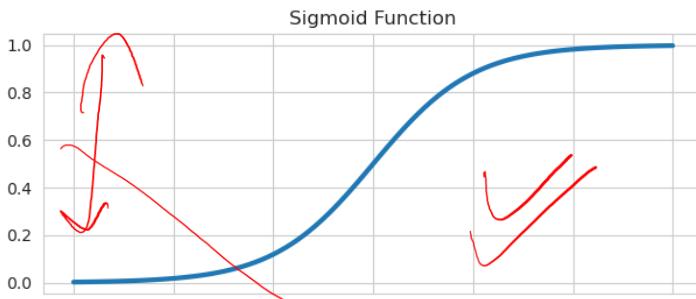
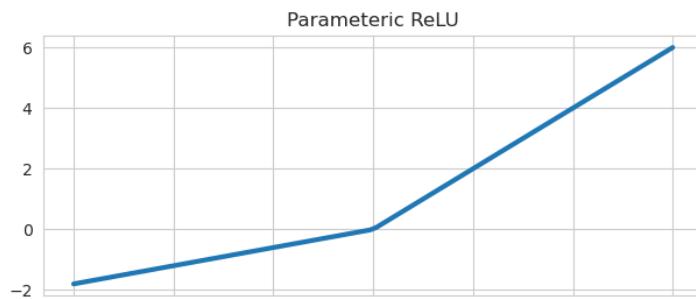
<https://playground.tensorflow.org/>



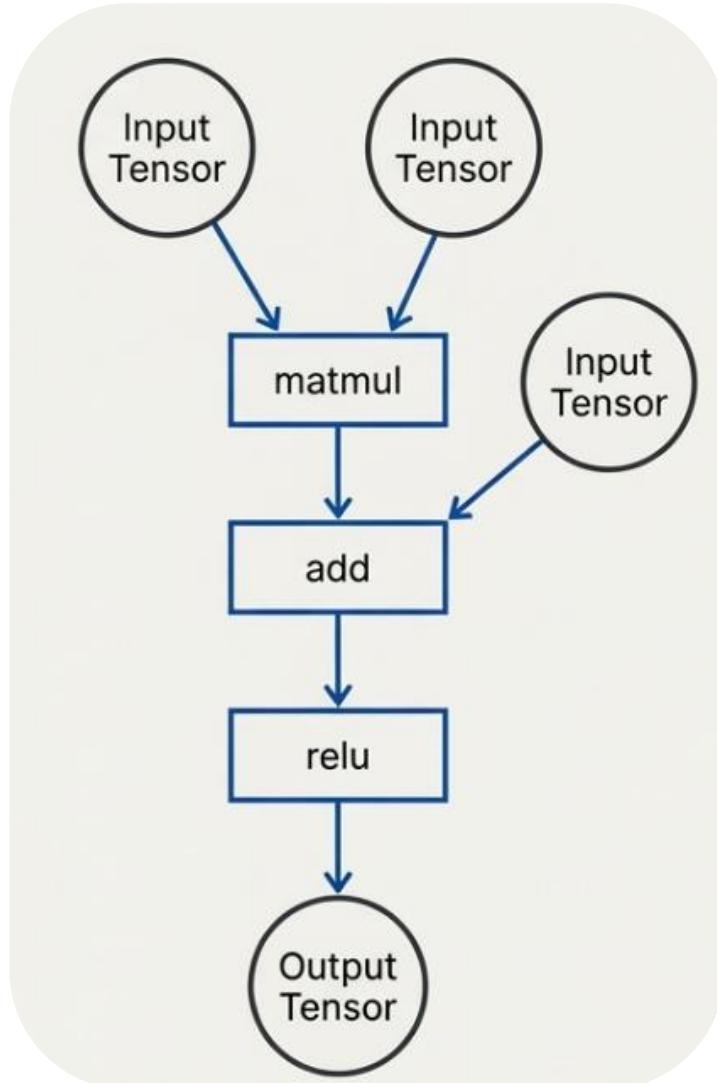
<https://tinyurl.com/dlframeworks>

<https://github.com/sakharamg/DeepLearningFrameworks>

Activation Functions

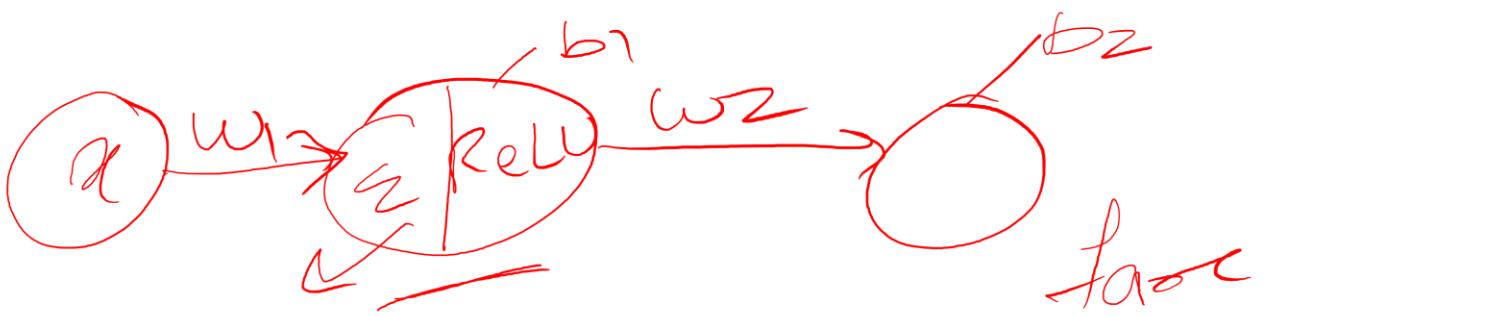


Recap: Computation Graph



- **Eager Execution:** ops run immediately (no “compile step”)
- **Debugging = inspection:** print **shape** / **dtype** / **device** early
- **Core Primitives:** **Tensor + Autograd** → everything else builds on this

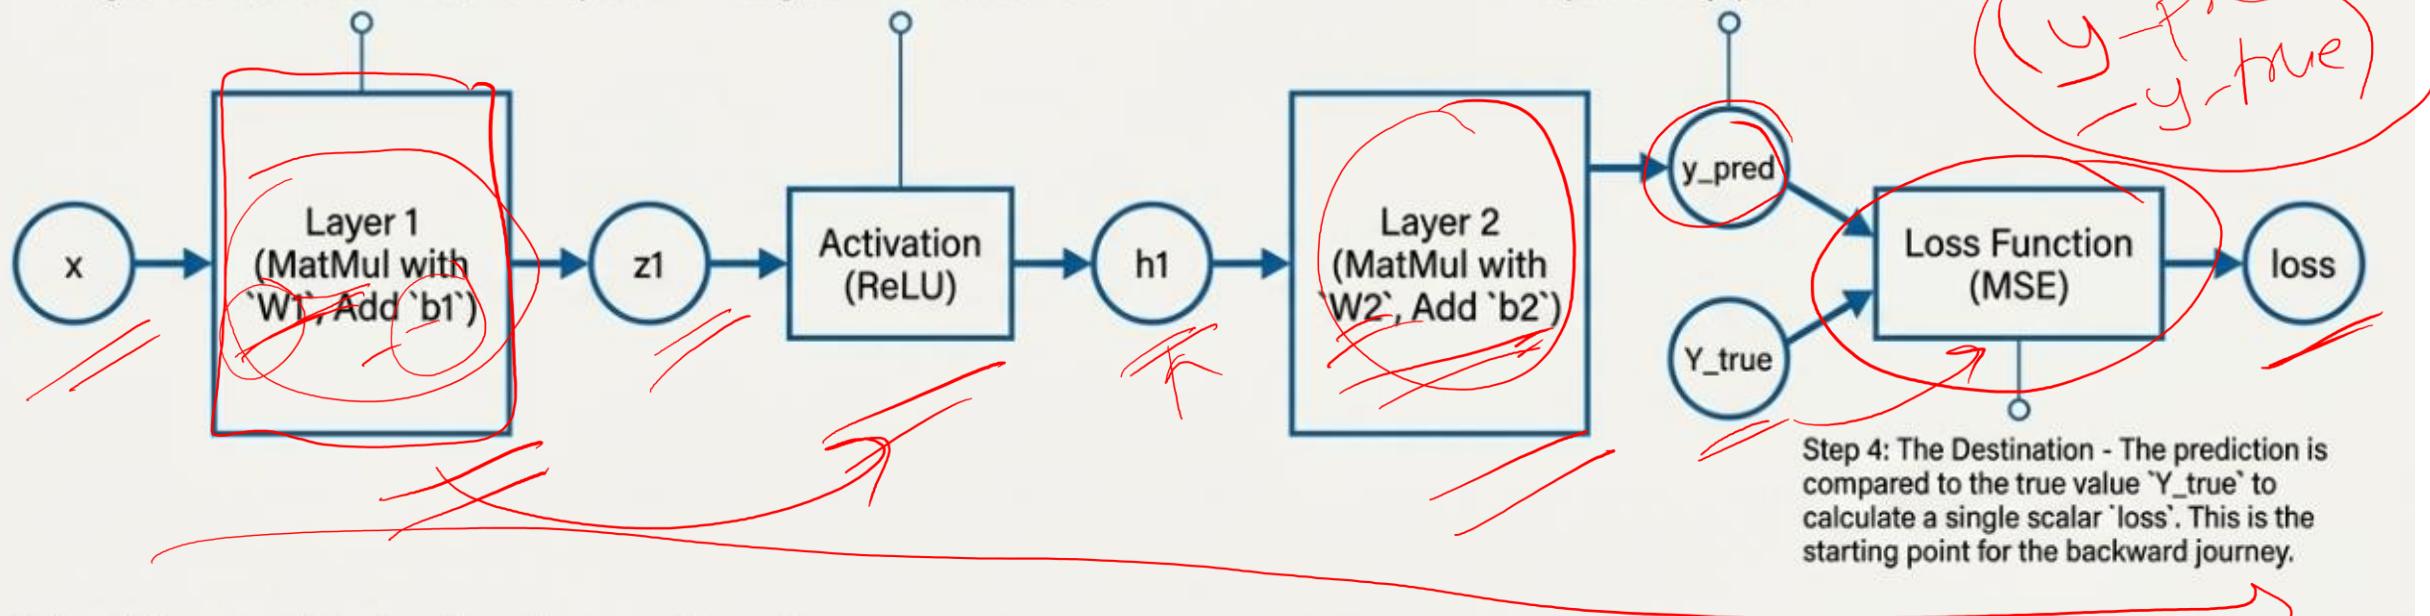
Forward Pass



Step 1: Linear Transformation - Input 'x' is transformed by weights 'W1' and bias 'b1'. Autograd records the 'AddmmBackward' operation.

Step 2: Non-Linear Activation - The result passes through a ReLU function. Autograd notes 'ReluBackward'.

Step 3: Final Output - The process repeats for the next layer to produce a prediction 'y_pred'.



What Happens During the Forward Pass?

The forward pass executes the model's operations sequentially to compute an output. Crucially, as this happens, PyTorch builds the dynamic computation graph in the background, linking each output tensor to its creating function via 'grad_fn'. No gradients are calculated yet, but the 'map' is now complete.

loss.backward()

The Signal is Born

After the forward pass, we have a single scalar value: the loss. The entire goal of training is to minimize this value. To do this, we need to know how each parameter (W , b) affects the loss. This is the gradient.

The chain rule allows us to compute this:

If $loss = f(y)$ and $y = g(W)$, then the gradient of the loss with respect to the weight W is:

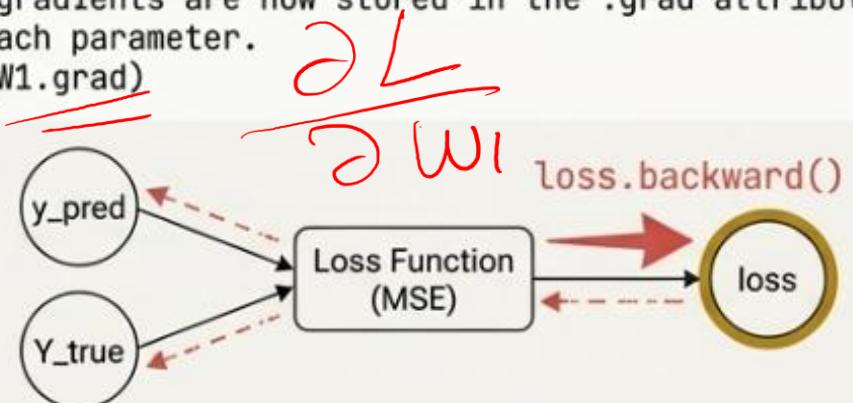
$$\frac{\partial loss}{\partial W} = \frac{\partial loss}{\partial y} * \frac{\partial y}{\partial W}.$$

Think of it as assigning blame. The chain rule mathematically distributes the total error back through each operation that contributed to it.

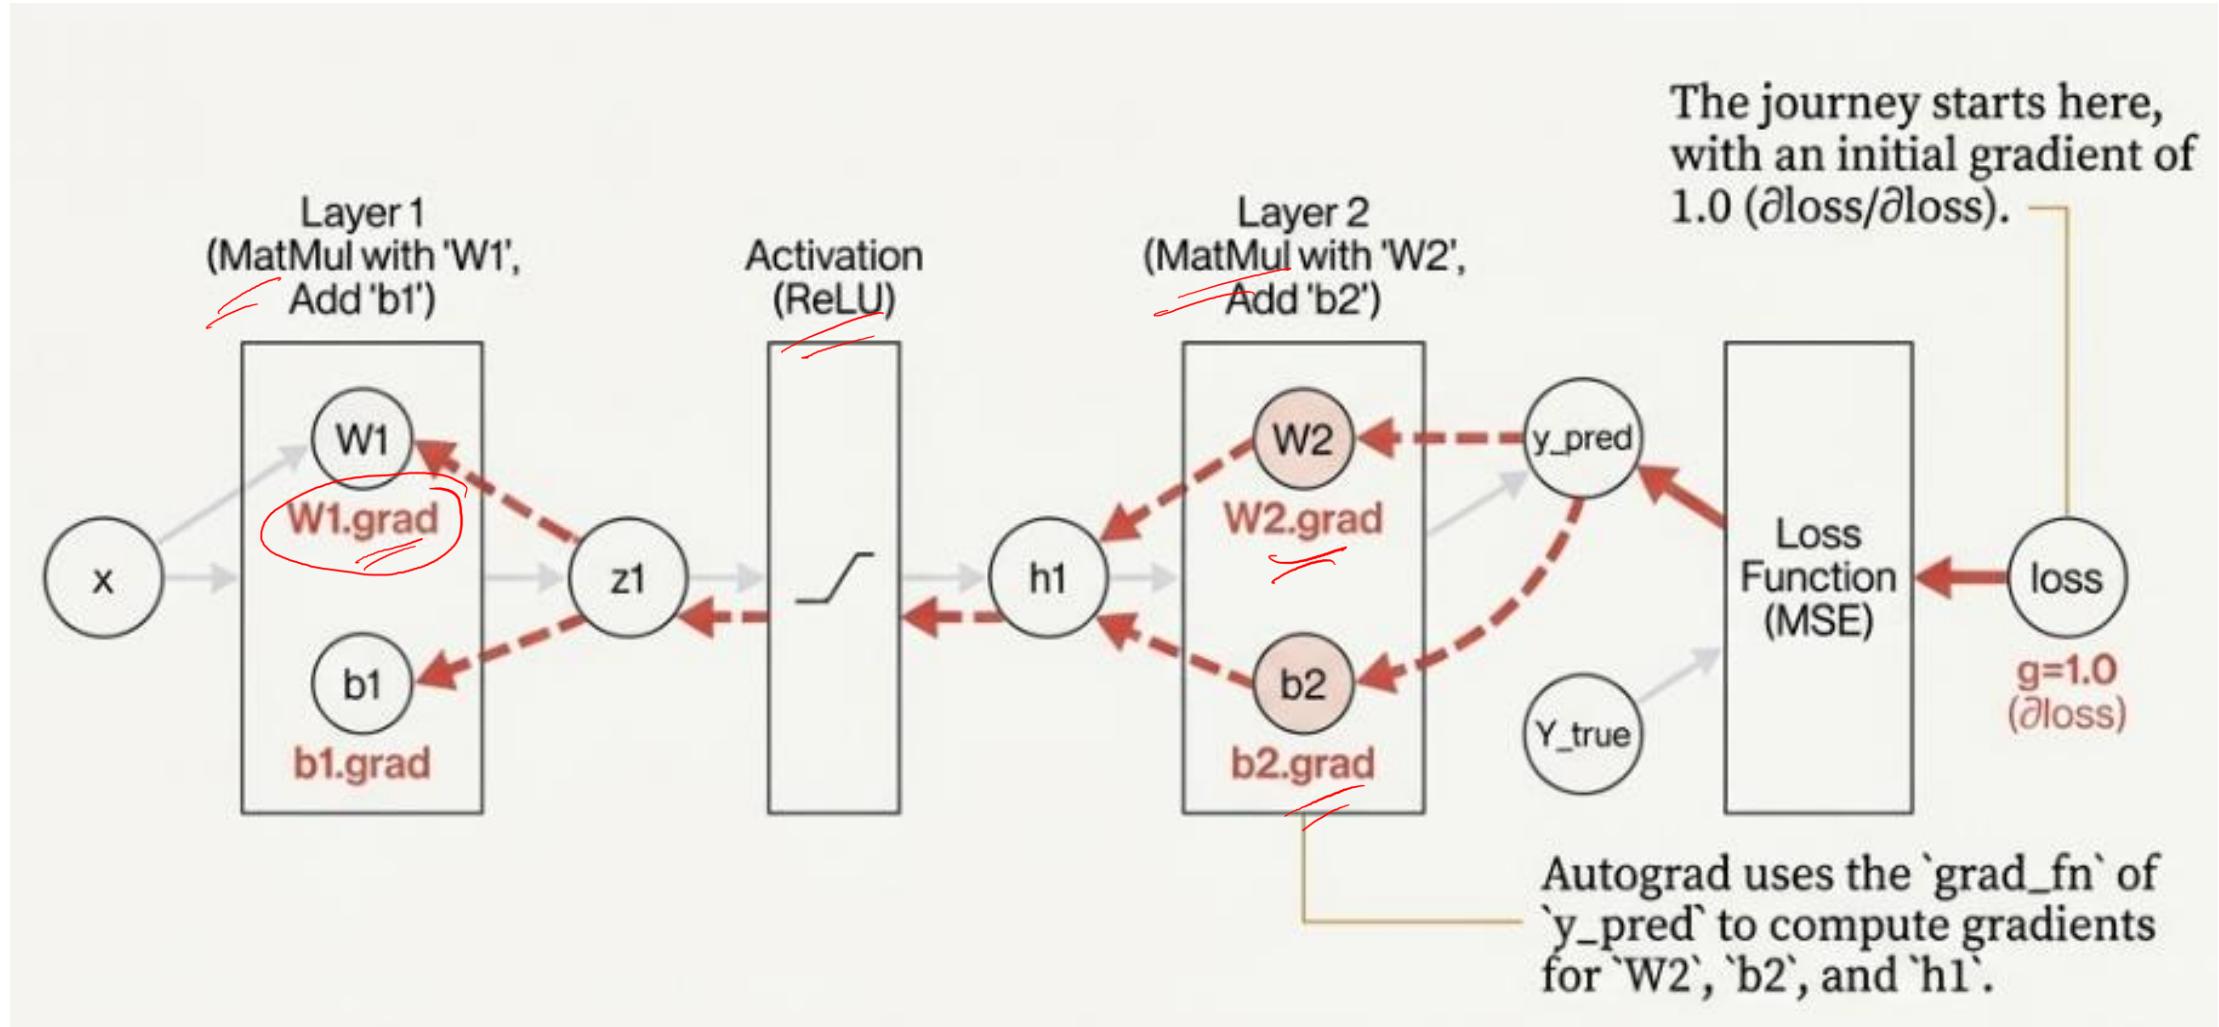
The Trigger: `loss.backward()`

In PyTorch, the entire backpropagation process is triggered by a single command on the final loss tensor.

```
# After computing the final loss...
loss = F.mse_loss(y_pred, Y_true)
# This one line starts the journey home for the gradients.
loss.backward()
# The gradients are now stored in the .grad attribute
# of each parameter.
print(W1.grad)
```

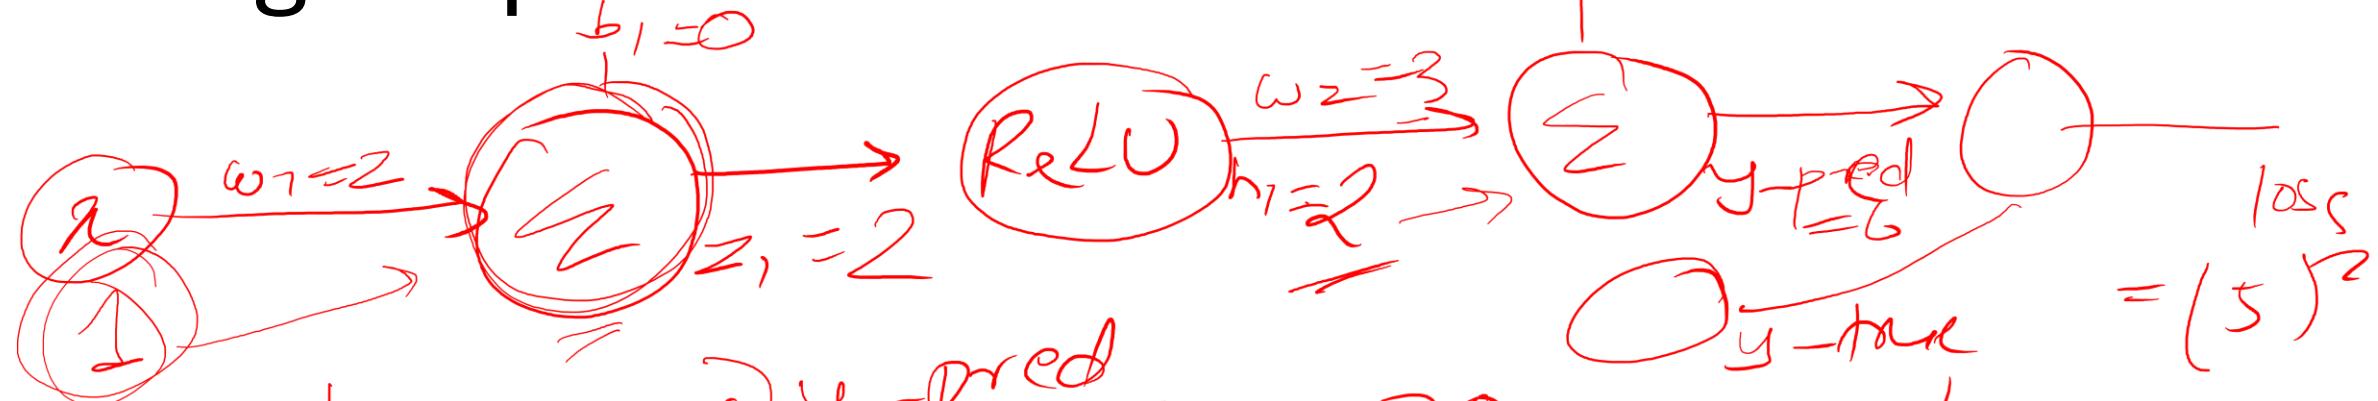


Back Propagation



$$\omega = \omega - \frac{\partial L}{\partial \omega}$$

Weight Updates



$$\frac{\partial L}{\partial w_2} = \frac{\partial L}{\partial y_{\text{pred}}} \cdot \frac{\partial y_{\text{pred}}}{\partial w_2} = 10 \cdot \frac{\partial y_{\text{pred}}}{\partial w_2} = 20$$

$$\frac{\partial L}{\partial b_2} = \frac{\partial L}{\partial y_{\text{pred}}} \cdot \frac{\partial y_{\text{pred}}}{\partial b_2} = 10 \cdot \frac{\partial y_{\text{pred}}}{\partial b_2} = 10$$

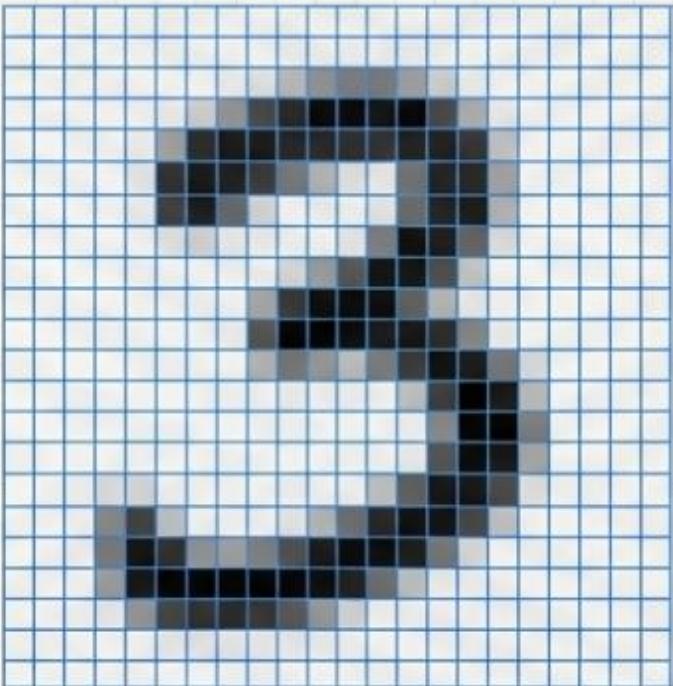
$$\frac{\partial L}{\partial w_1} = \frac{\partial L}{\partial y_{\text{pred}}} \cdot \frac{\partial y_{\text{pred}}}{\partial h_1} \cdot \frac{\partial h_1}{\partial z_1} \cdot \frac{\partial z_1}{\partial w_1} = 10 \cdot 3 = 30$$

$$= \frac{60}{30}$$

Lab

MNIST

INPUT



A 28x28 pixel grayscale image.

Tensor Shape: `1 x 28 x 28`

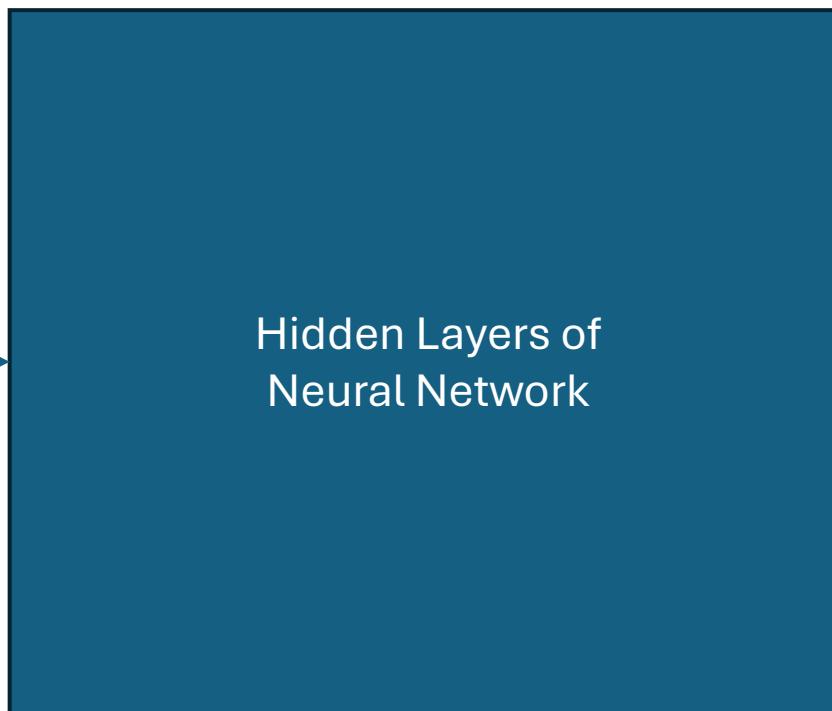
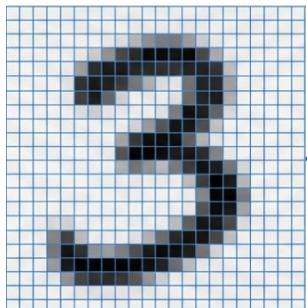
OUTPUT

0
1
2
3
4
5
6
7
8
9

A prediction of the digit's class.

Classes: `0, 1, 2, 3, 4, 5, 6, 7, 8, 9`

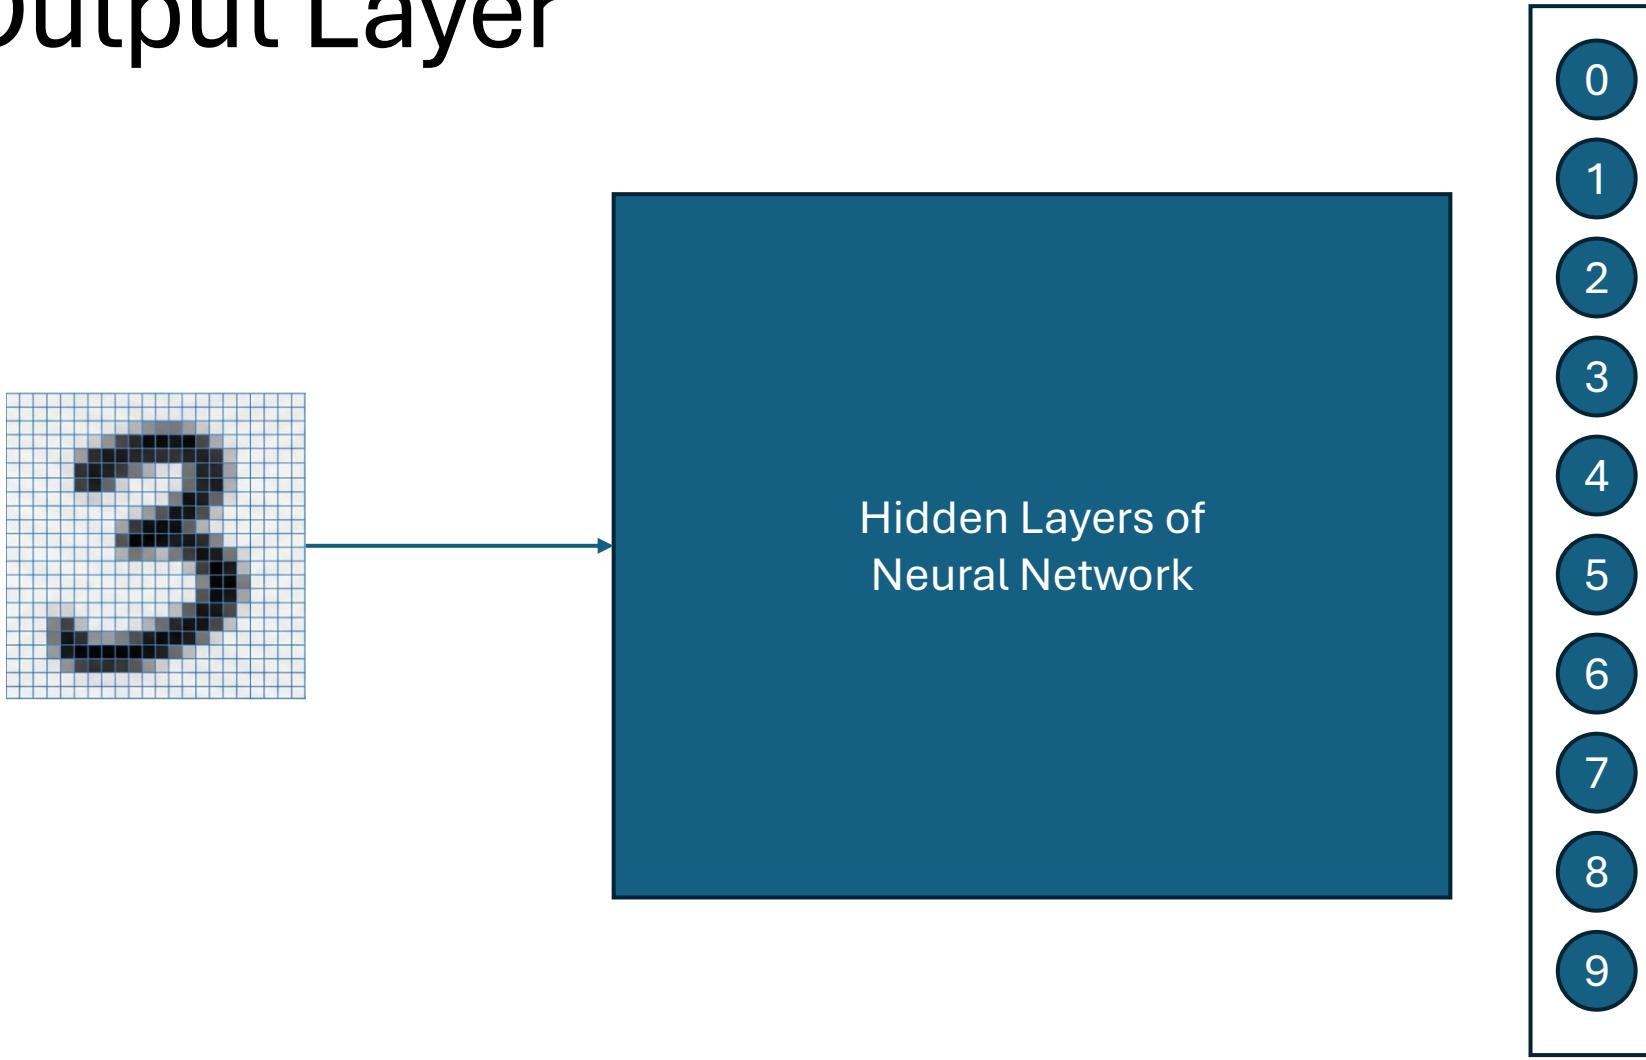
Output Layer



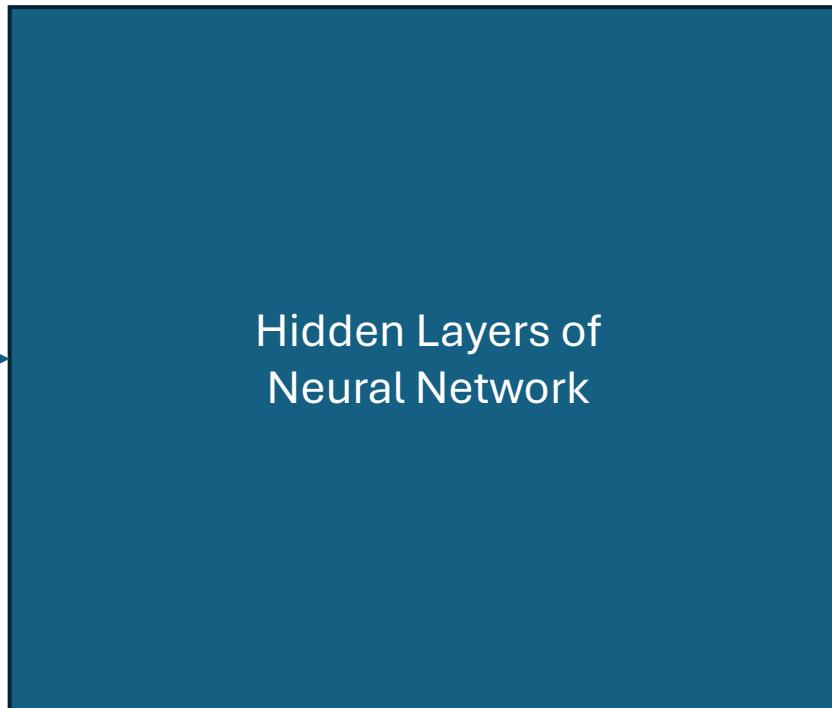
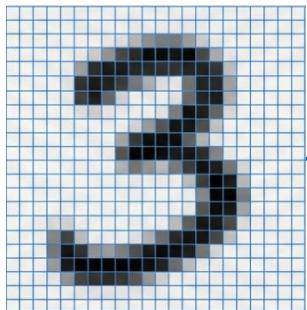
<https://tinyurl.com/dlframeworks>

<https://github.com/sakharamg/DeepLearningFrameworks>

Output Layer



Output Layer

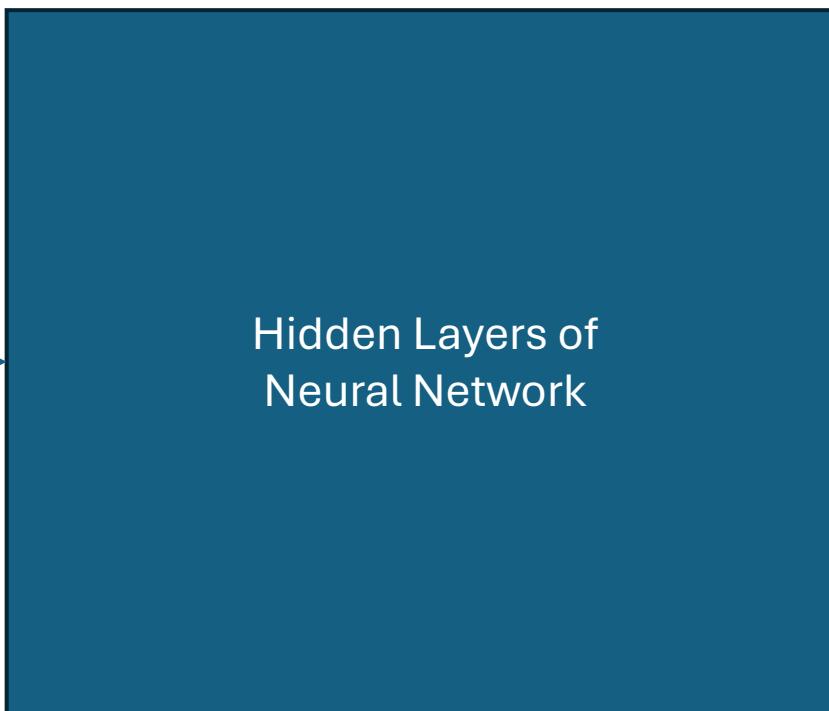
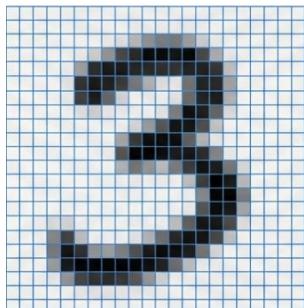


Hidden Layers of
Neural Network

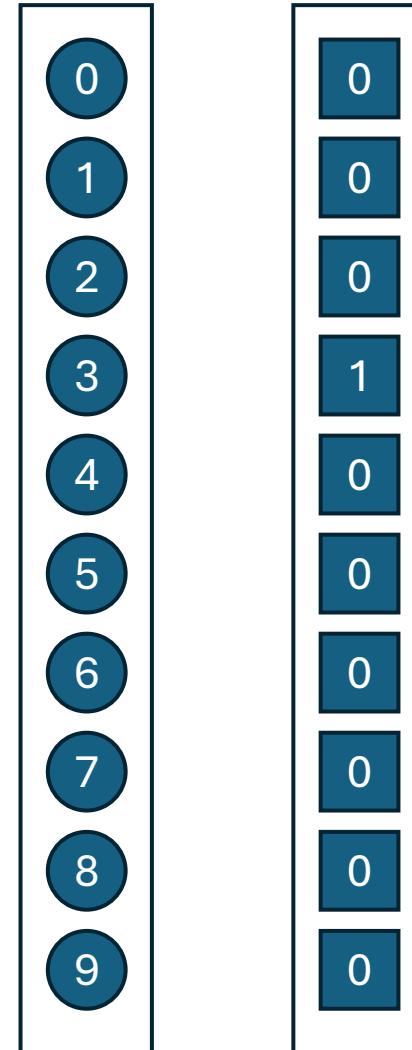


`nn.Linear(__, 10)`

Output Layer

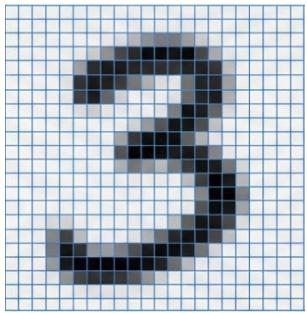


Hidden Layers of
Neural Network



`nn.Linear(__, 10)`

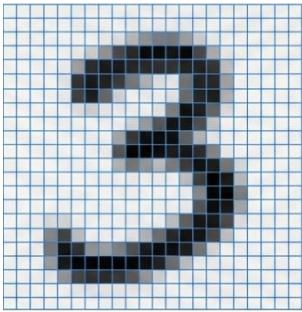
Softmax



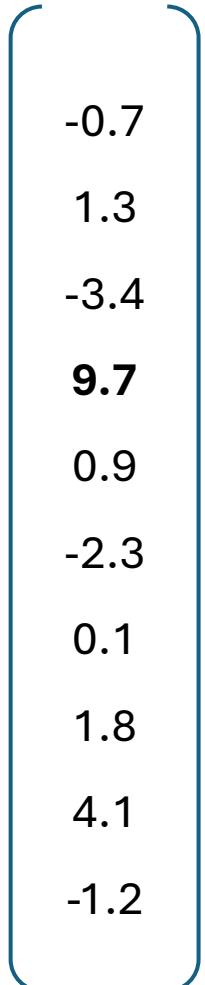
<https://tinyurl.com/dlframeworks>

<https://github.com/sakharamg/DeepLearningFrameworks>

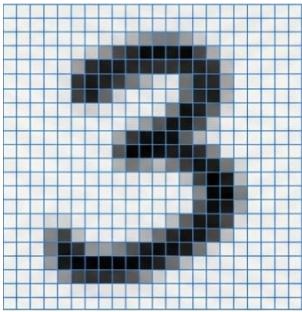
Softmax



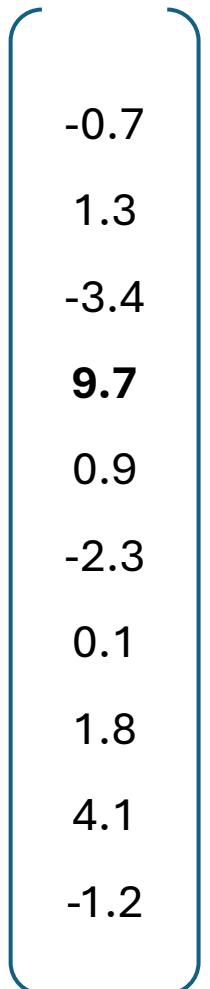
Output Layer



Softmax



Output Layer



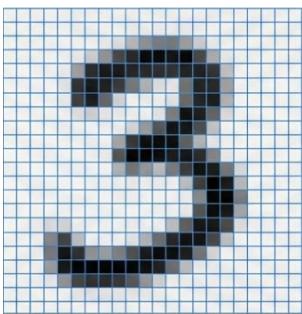
Softmax
activation function

$$\frac{e^{z_i}}{\sum_{j=1}^K e^{z_j}}$$

<https://tinyurl.com/dlframeworks>

<https://github.com/sakharamg/DeepLearningFrameworks>

Softmax



Output Layer

-0.7
1.3
-3.4
9.7
0.9
-2.3
0.1
1.8
4.1
-1.2

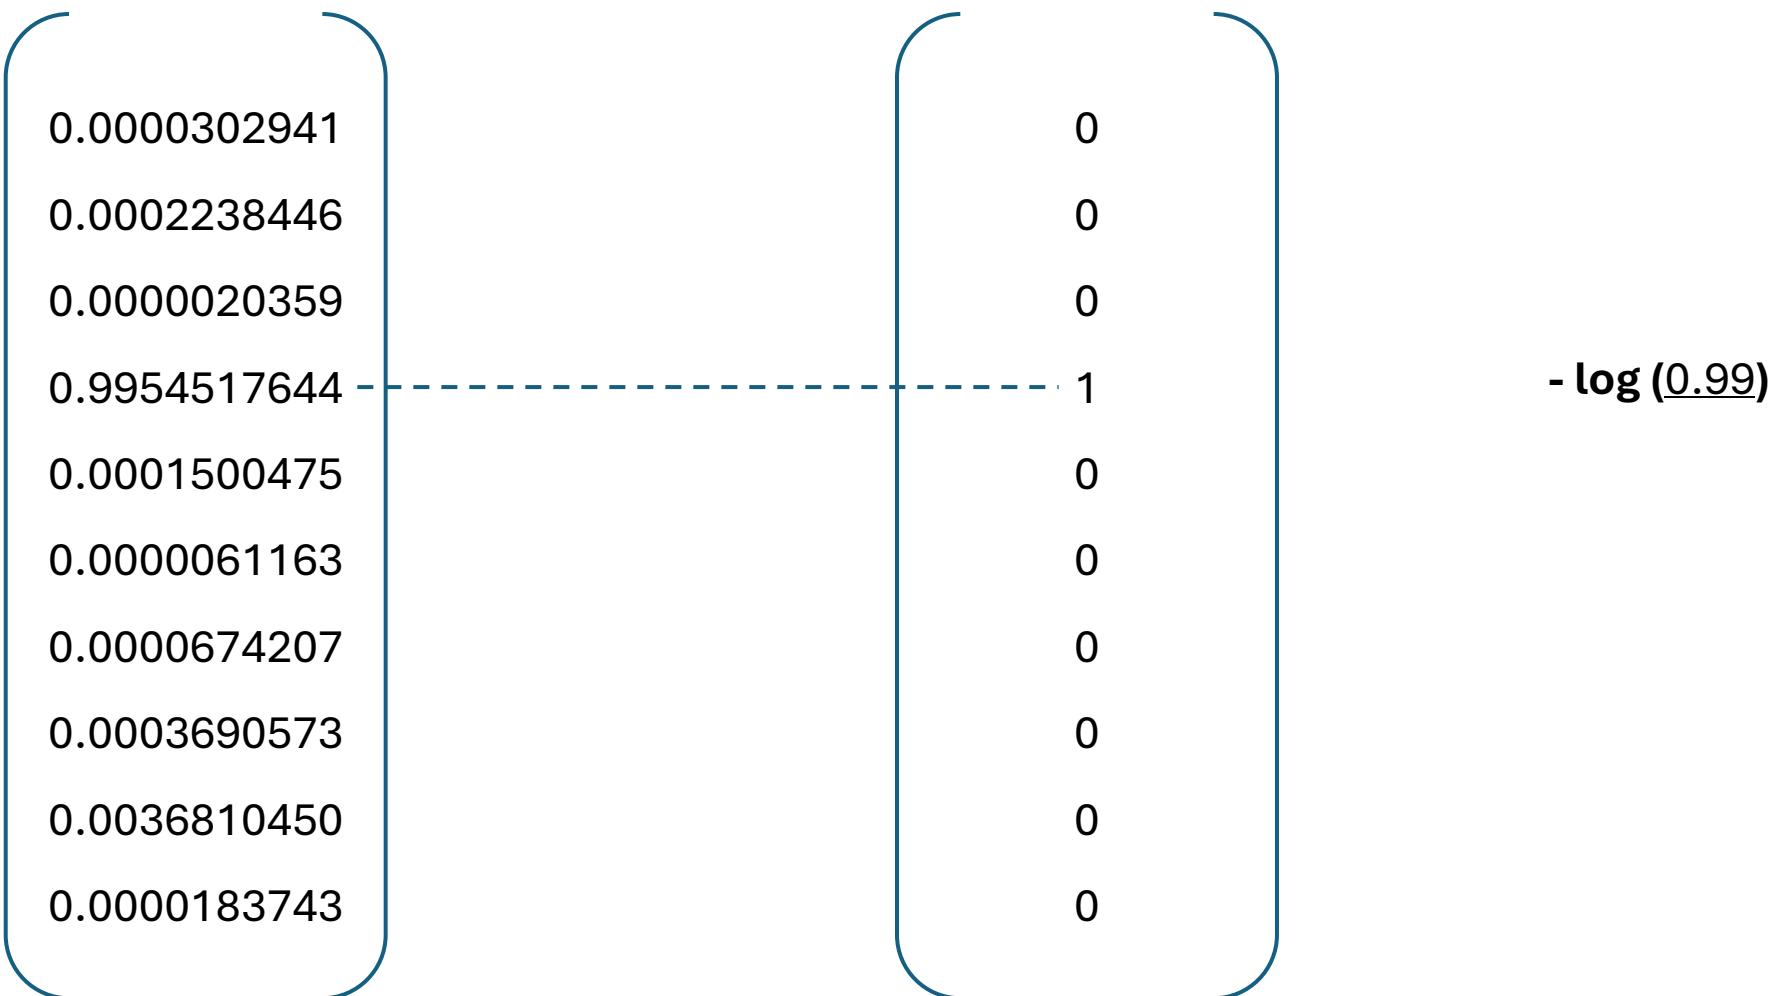
Logits

Softmax
activation function

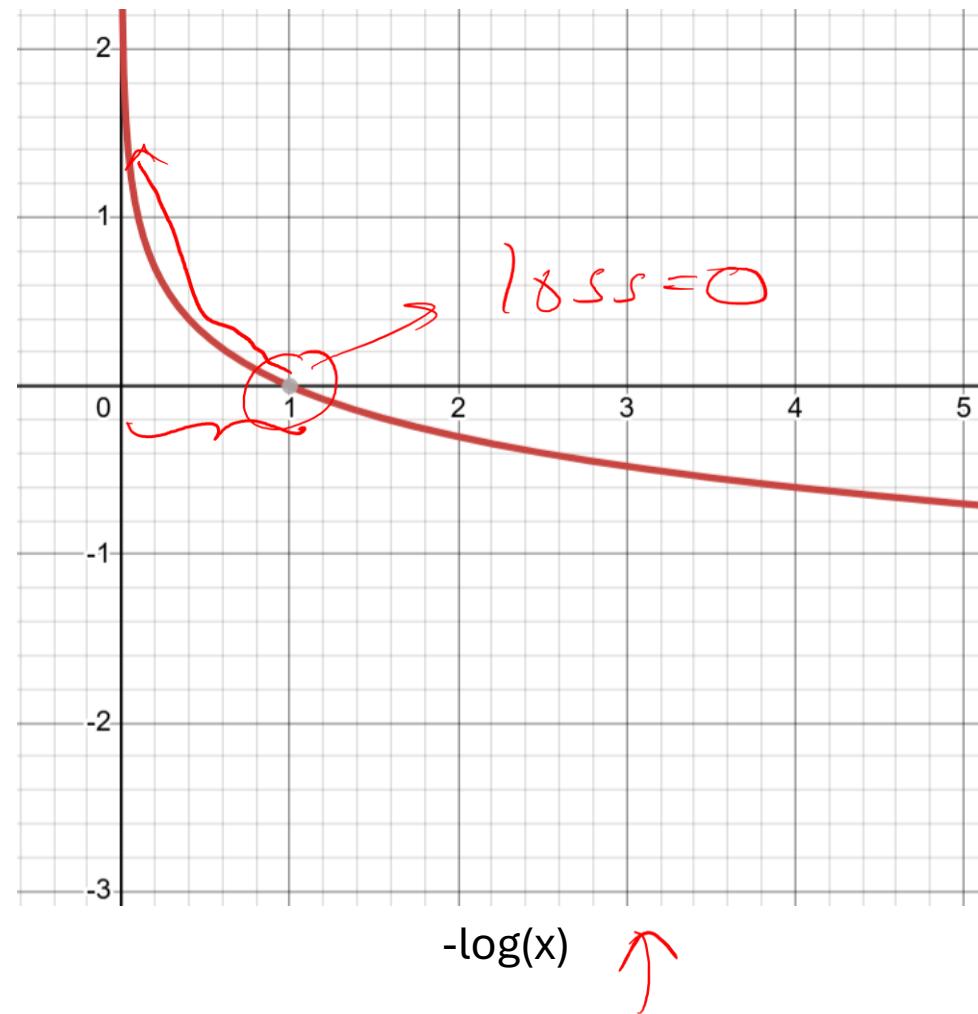
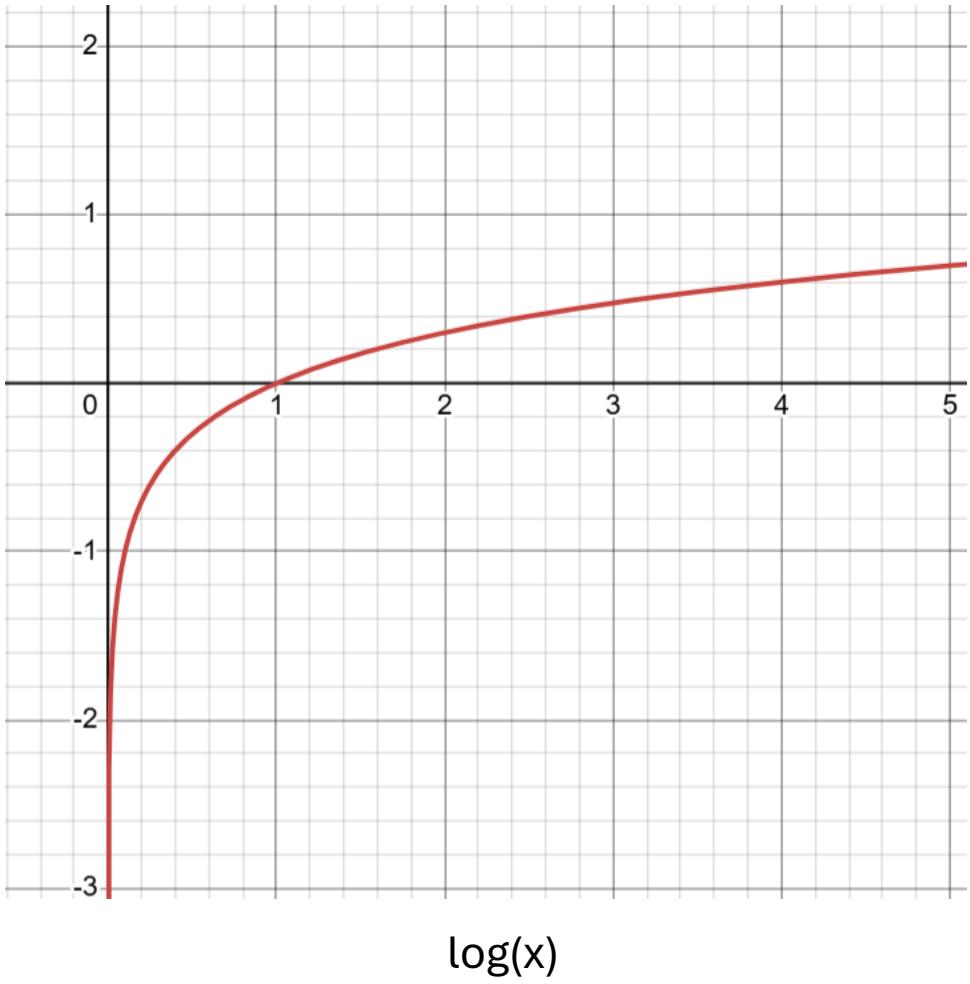
$$\frac{e^{z_i}}{\sum_{j=1}^K e^{z_j}}$$

0.0000302941
0.0002238446
0.0000020359
0.9954517644
0.0001500475
0.0000061163
0.0000674207
0.0003690573
0.0036810450
0.0000183743

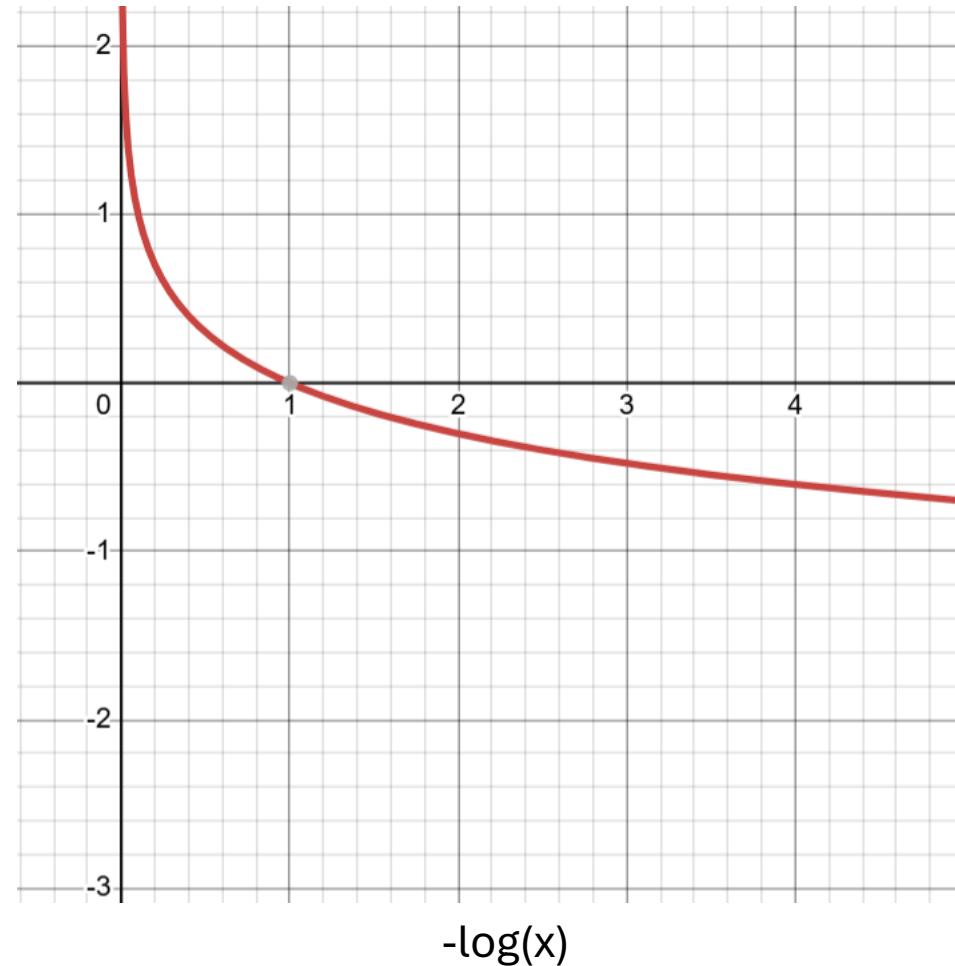
Negative Log Likelihood



Negative Log Likelihood



Negative Log Likelihood



Lesser Confidence \rightarrow More Penalty



Cross Entropy Loss

logsumexp

1. $P = \text{log_softmax}(\text{logits})$... Supress Probabilities
2. $\text{NLL}(\log_P) = -\log(P_{\text{true_class}})$

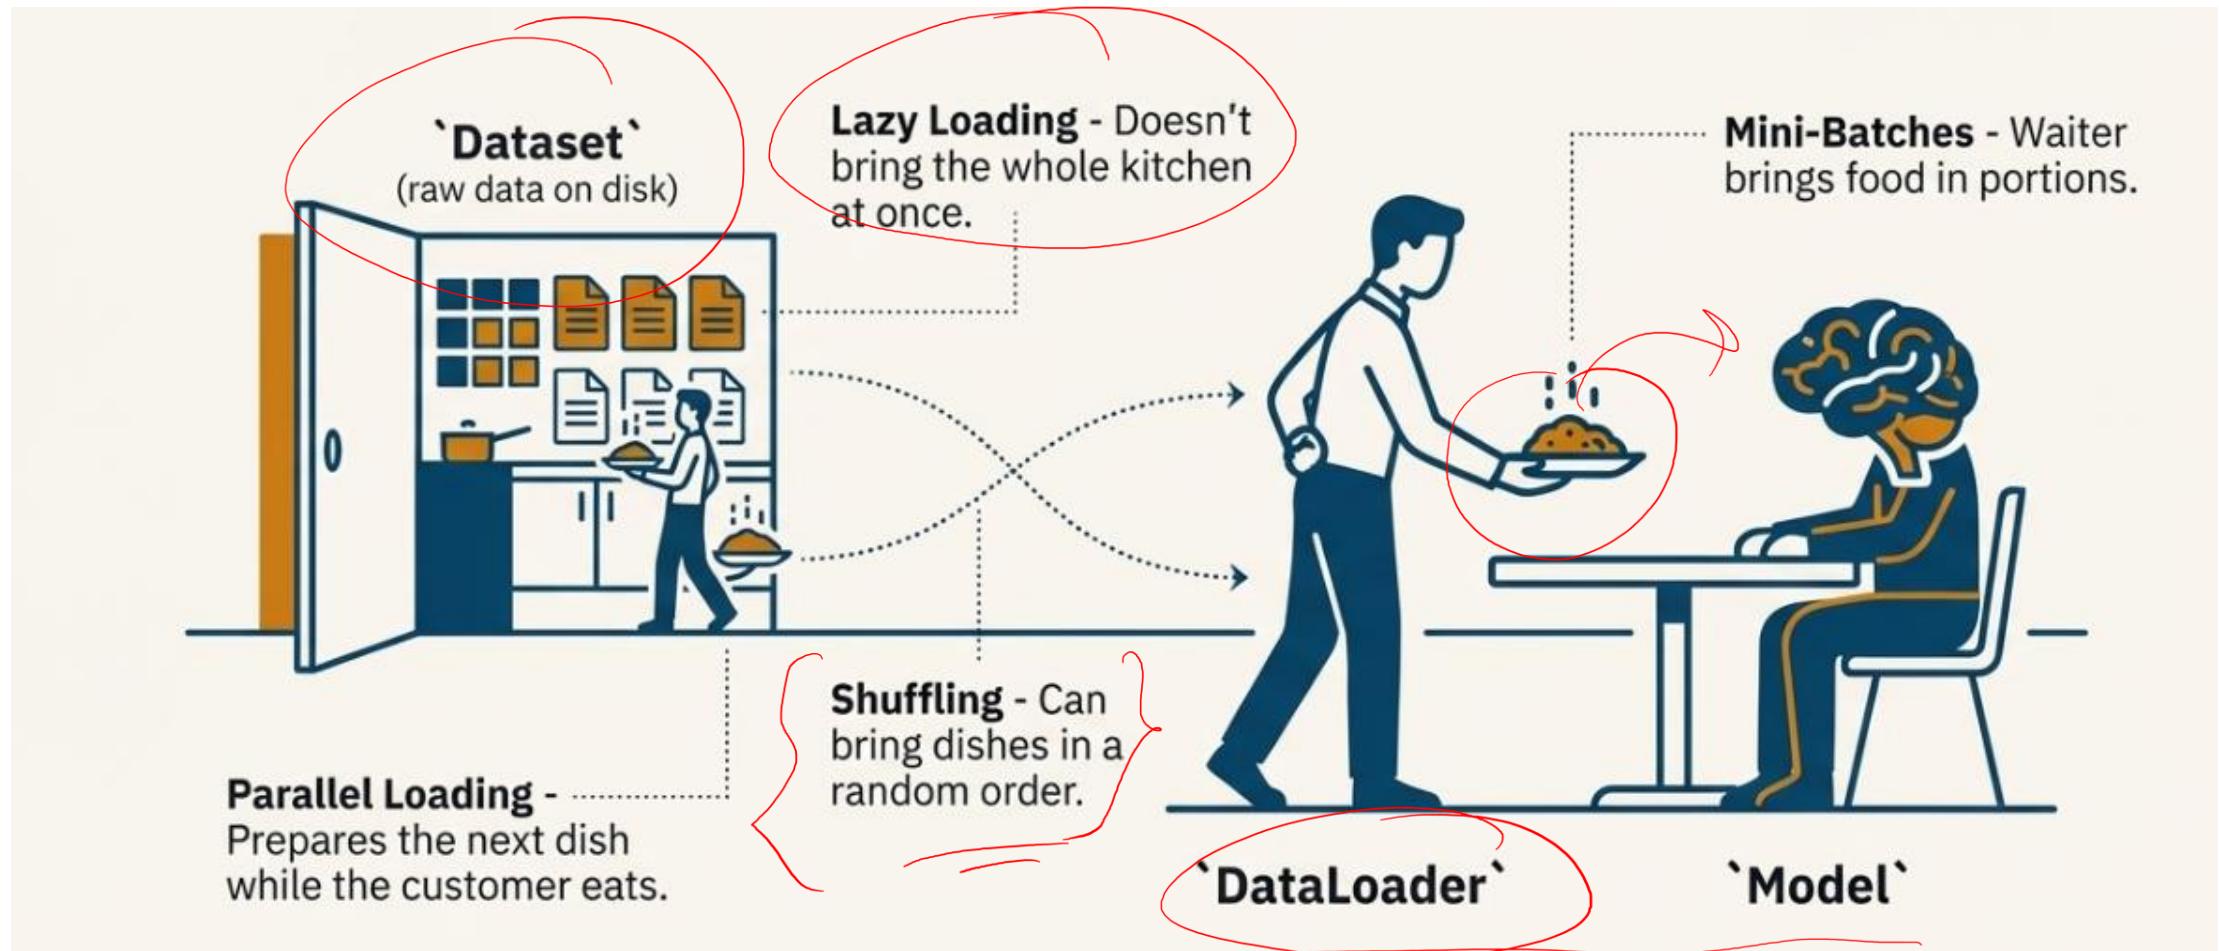
Softmax activation function

$$\log \frac{e^{z_i}}{\sum_{j=1}^K e^{z_j}}$$

$$\log \text{softmax}(z)_k = z_k - \log \sum_j e^{z_j}$$

- if logits are large (e.g., 80, 100), e^z can overflow to inf
- if logits are very negative (e.g., -80), e^z underflows to 0
- then $\log(\text{softmax})$ becomes $\log(0)$ or $\log(\text{inf}) \rightarrow -\text{inf}/\text{inf}$, NaNs, unstable gradients

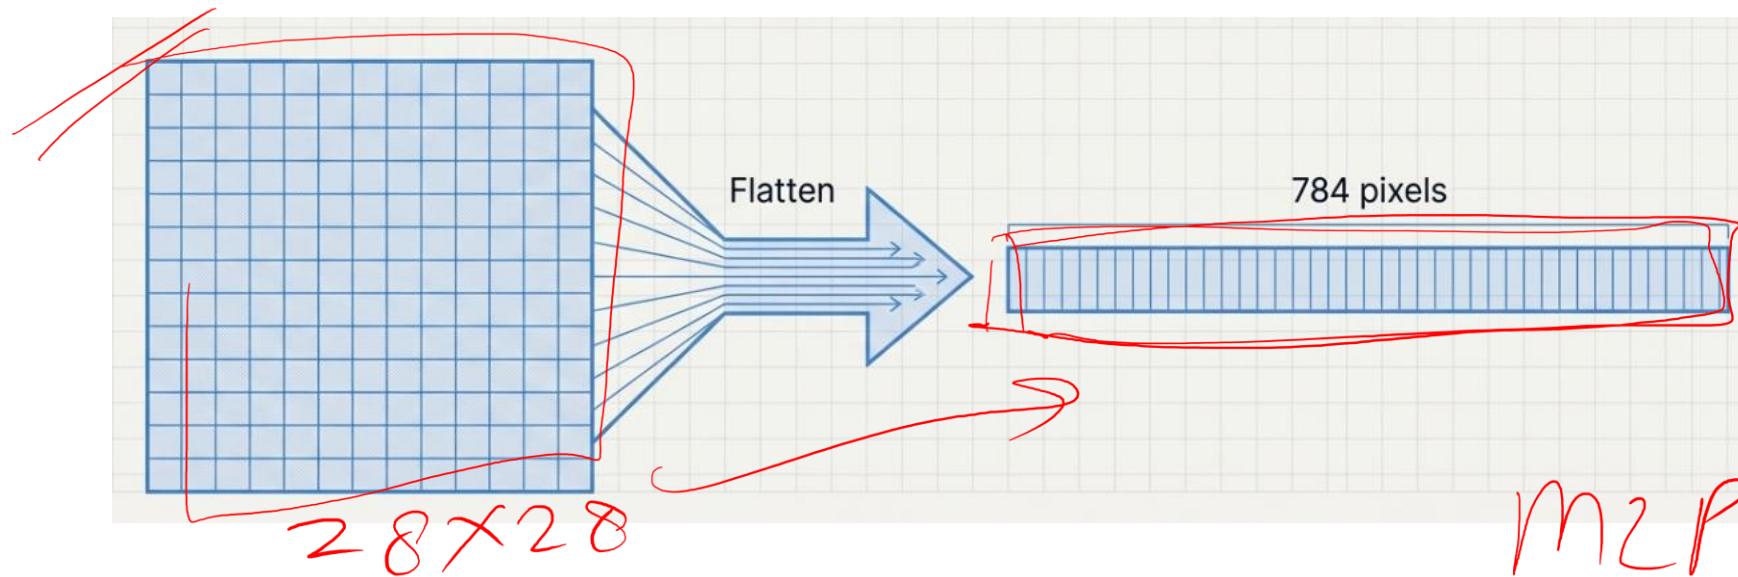
DataLoader



<https://tinyurl.com/dlframeworks>

<https://github.com/sakharamg/DeepLearningFrameworks>

Flatten

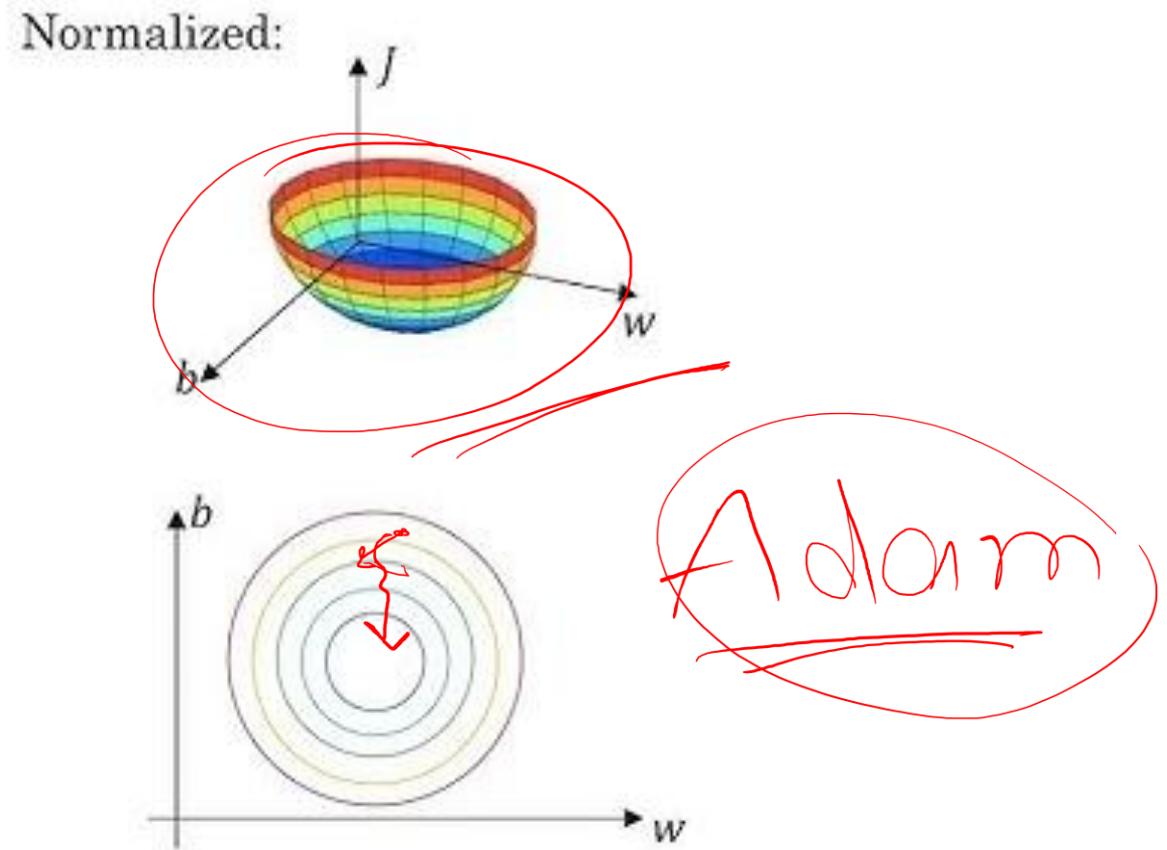
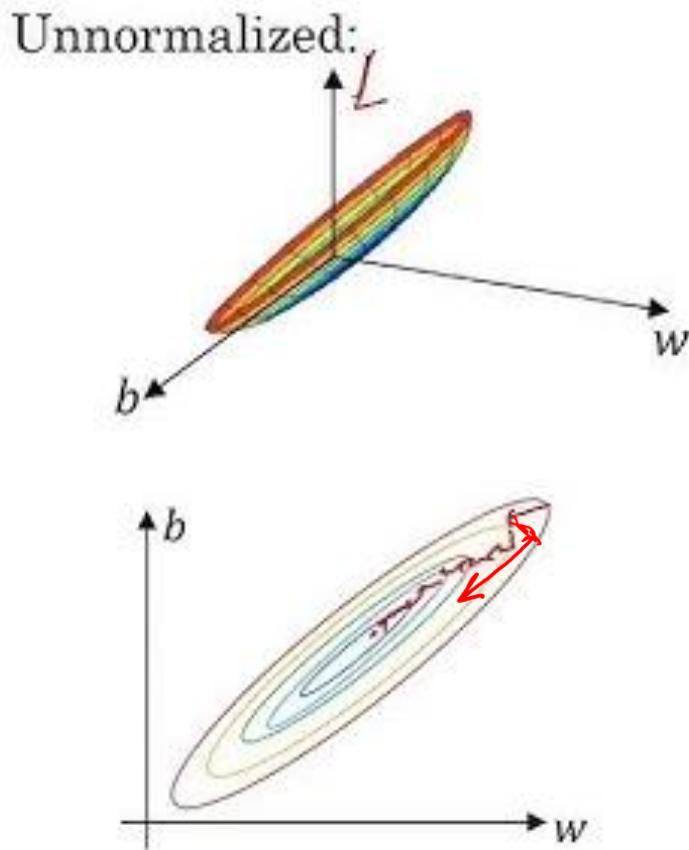


A 28x28 images contains 784 pixels.

MLP cannot process 2D spatial data directly. It requires a 1D vector as input.

The first step in our model's logic must be to **Flatten** the $(N, 1, 28, 28)$ image tensor into a $(N, 784)$ vector.

Normalize()



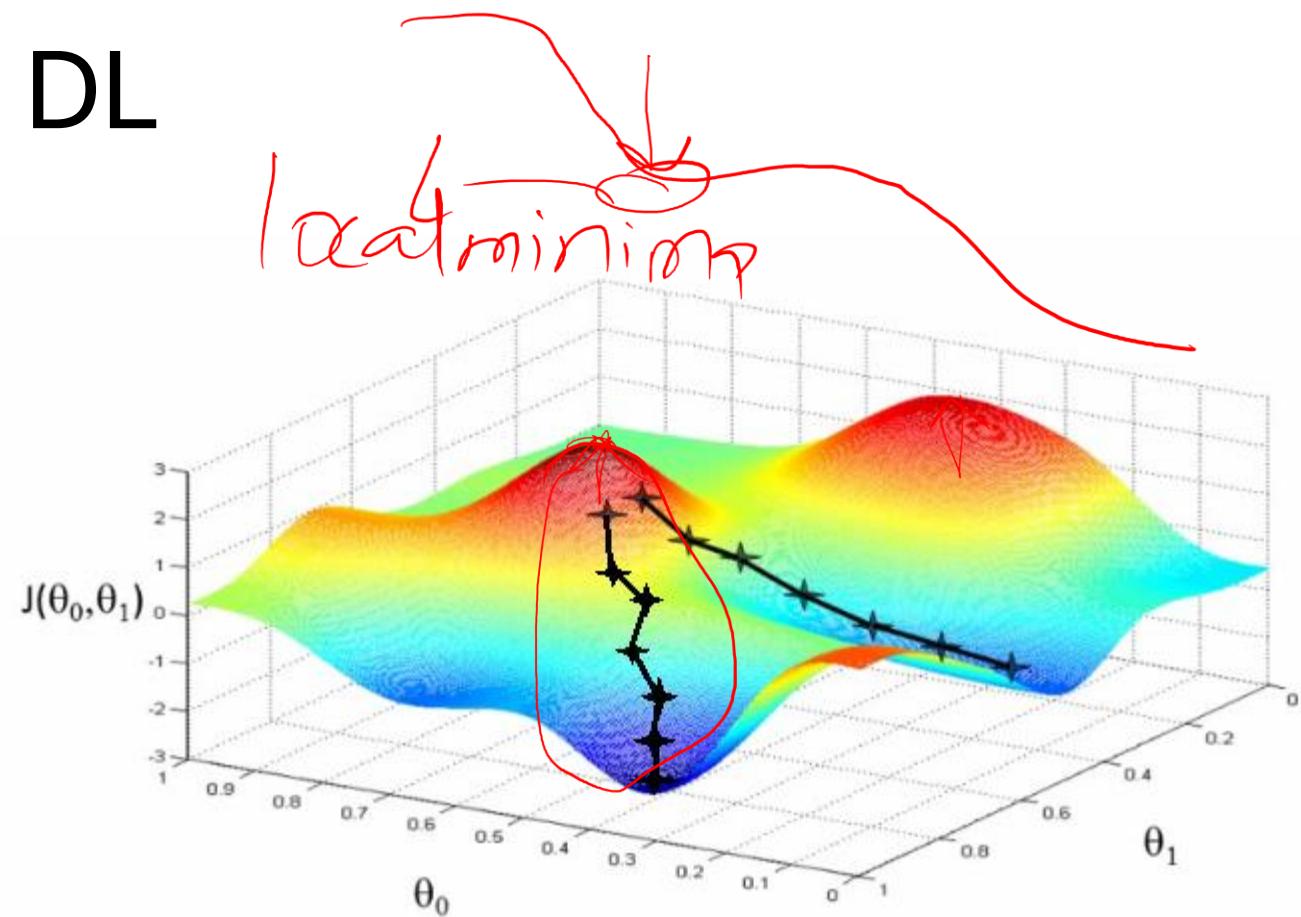
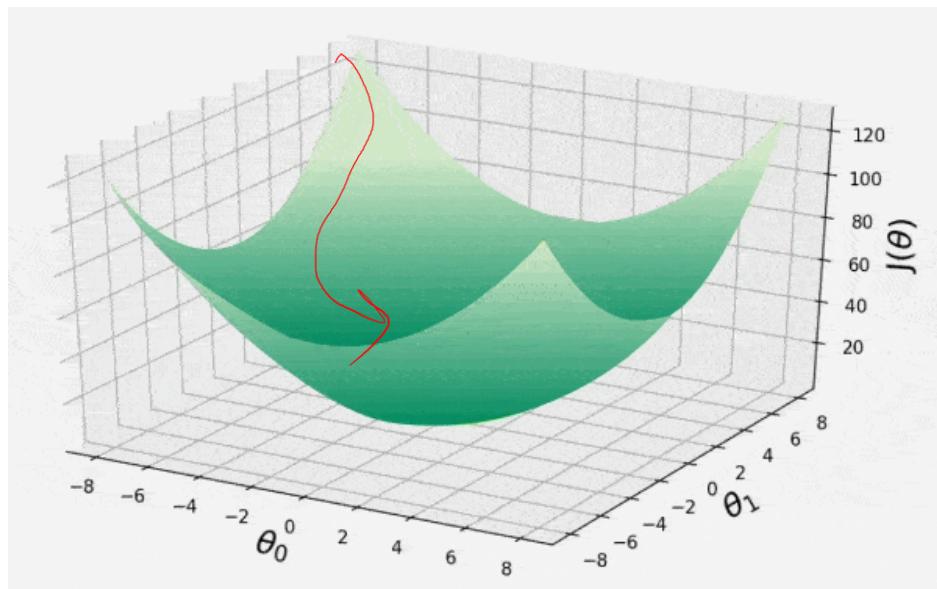
Source: DeepLearningAI (C2W1L09)

Andrew Ng

<https://tinyurl.com/dlframeworks>

<https://github.com/sakharamg/DeepLearningFrameworks>

Gradient Descend - DL



<https://tinyurl.com/dlframeworks>

<https://github.com/sakharamg/DeepLearningFrameworks>

Lab

<https://tinyurl.com/dlframeworks>

<https://github.com/sakharamg/DeepLearningFrameworks>

Regularization: Weight Decay and Dropout

Weight Decay: L2 Regularization

Dropout: Drop fraction of neurons

Weight Decay

$$L_{total} = L_{data} + \lambda \sum w^2$$

Mechanism

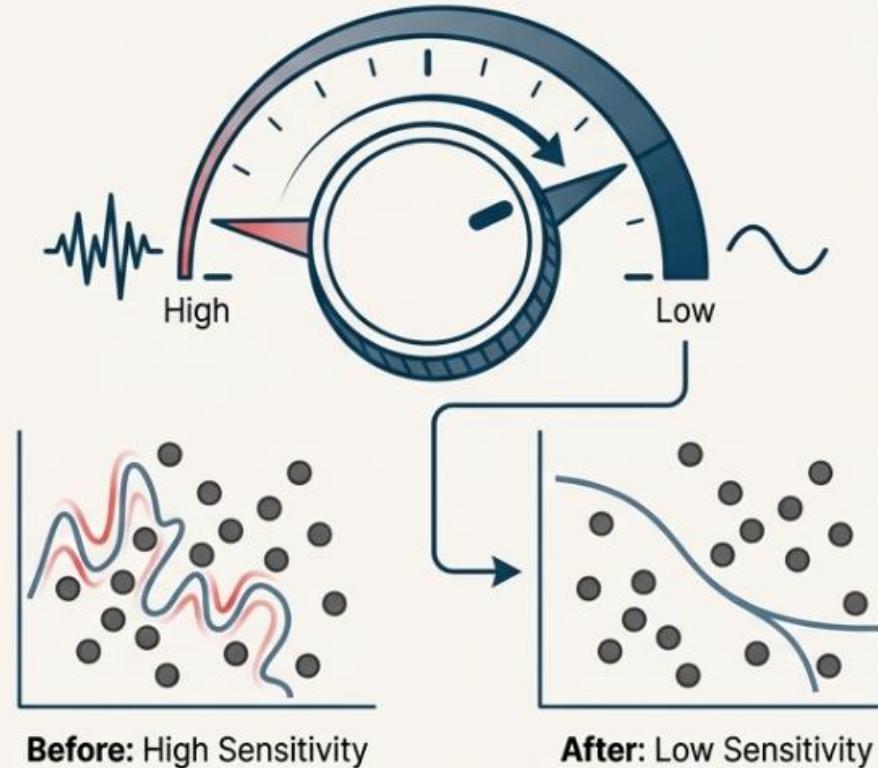
Adds a penalty term to the loss function proportional to the square of the weights ($\lambda \sum w^2$). During training, this continuously pulls weights towards zero, making large weights “expensive.”

Problem Solved

Prevents overly sharp and sensitive models. Large weights mean small input changes can cause large output changes.

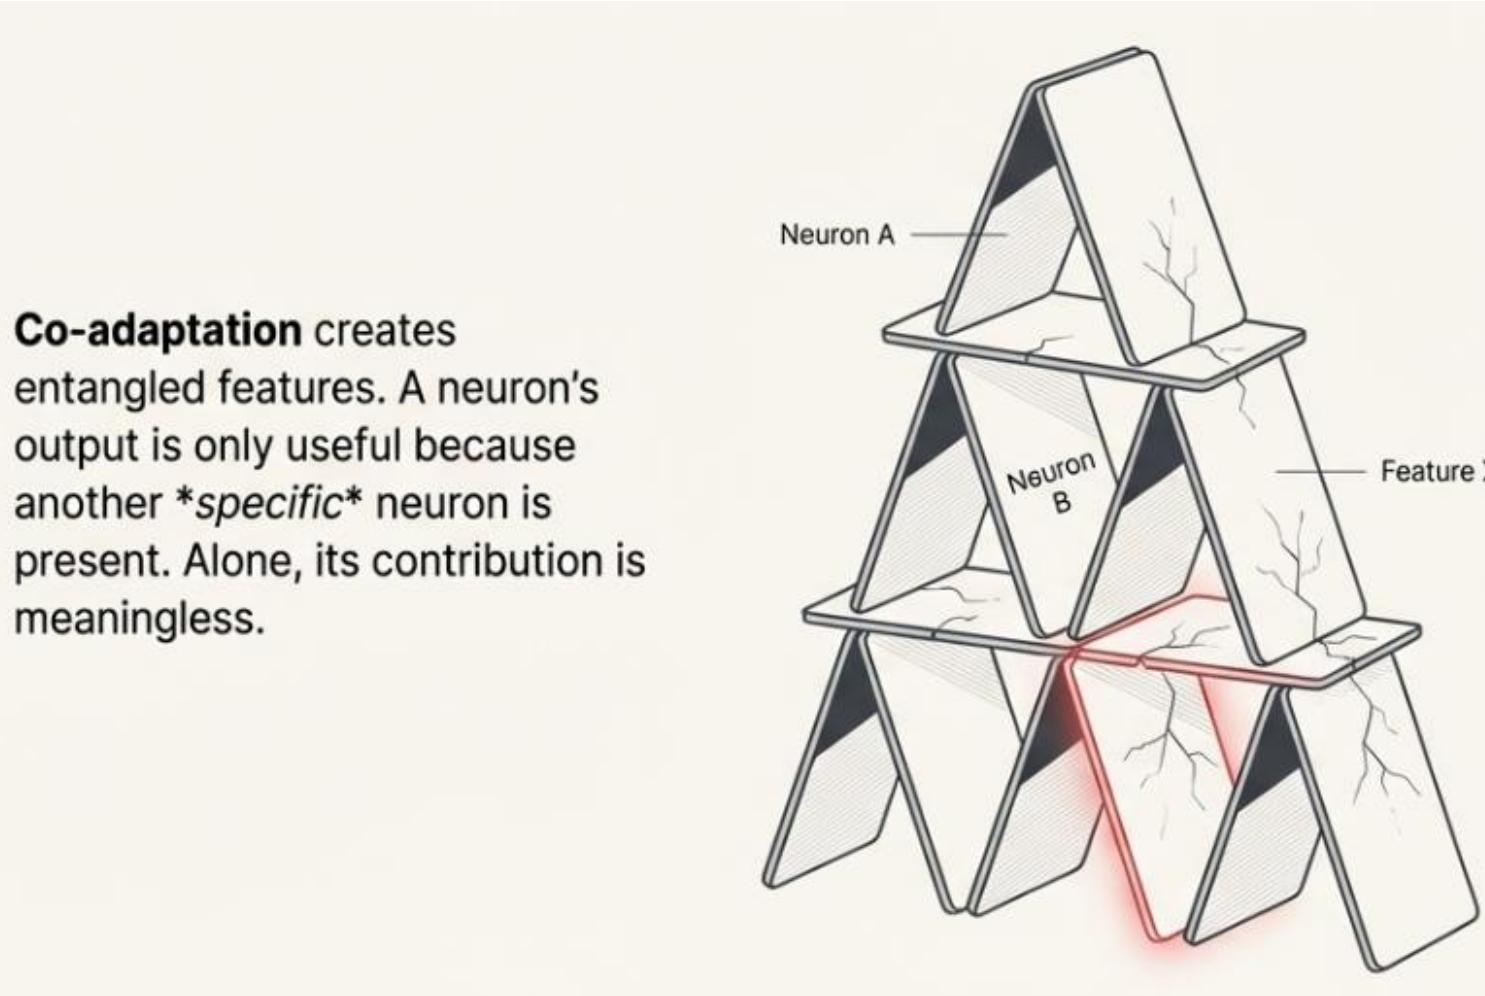
Intuition Analogy

Think of a microphone’s gain. High gain is sensitive and amplifies noise. Lowering the gain produces a clearer, more stable signal.



“Use many small contributions, not a few dominant ones.”

Co-Adaptation

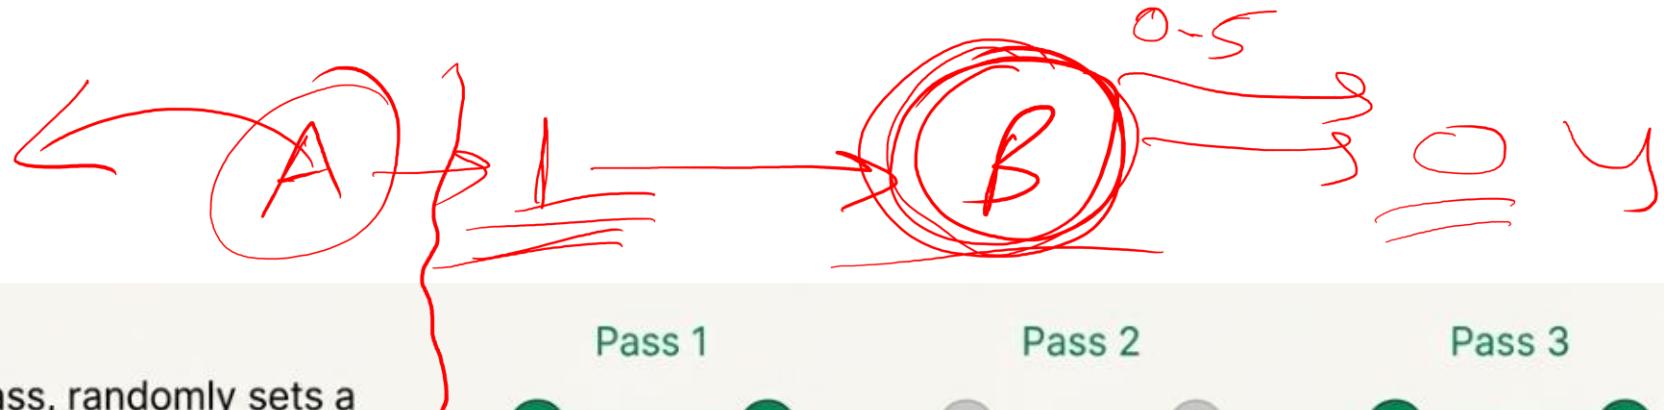


Am I a Siberian Husky?

<https://tinyurl.com/dlframeworks>

<https://github.com/sakharamg/DeepLearningFrameworks>

Dropout



Mechanism

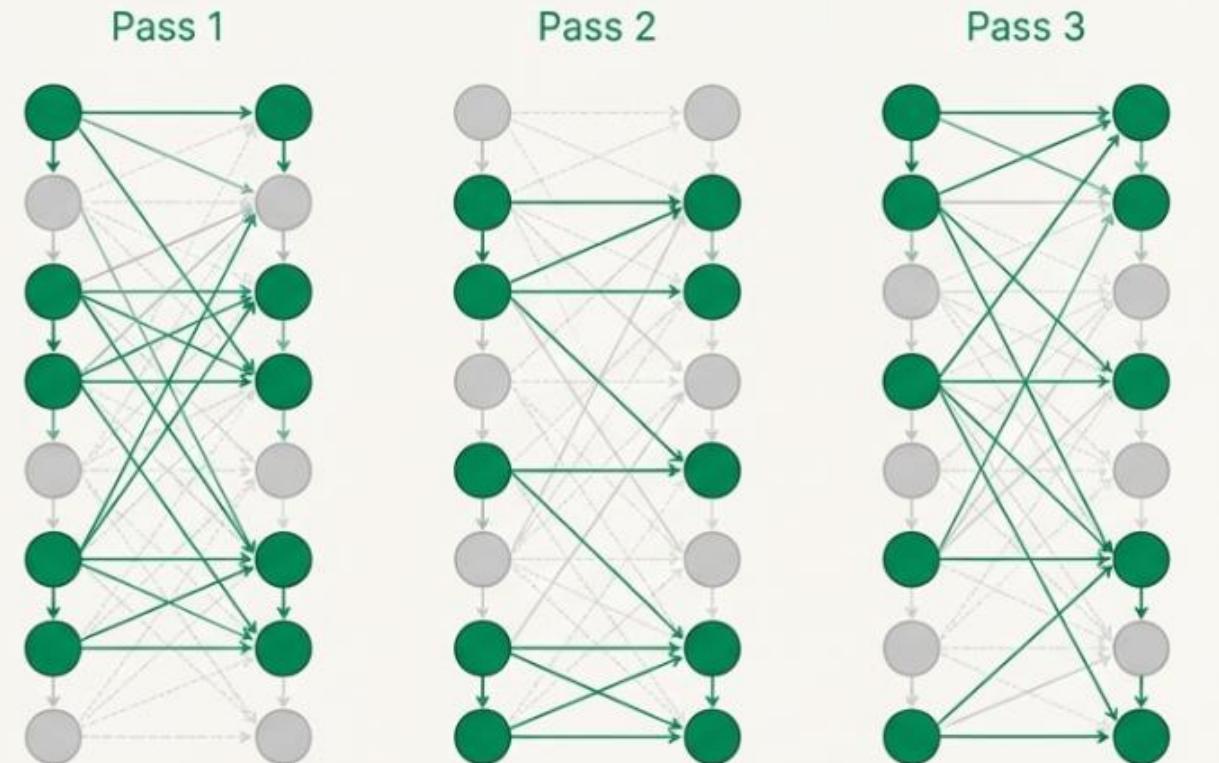
During each training pass, randomly sets a fraction of neuron activations to zero. This forces the network to function even with missing components. The network changes with every batch.

Problem Solved

Directly targets fragile feature dependencies (co-adaptation).

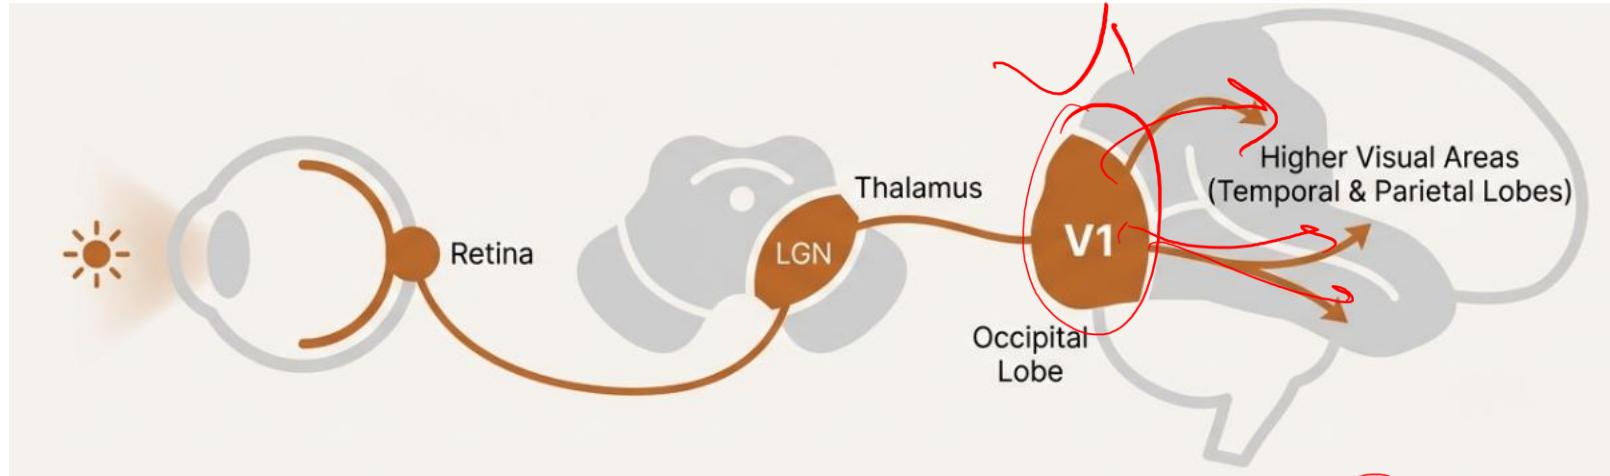
Intuition Analogy

A team project where any member might be absent on any given day. To succeed, everyone must learn the full context instead of relying on one person.

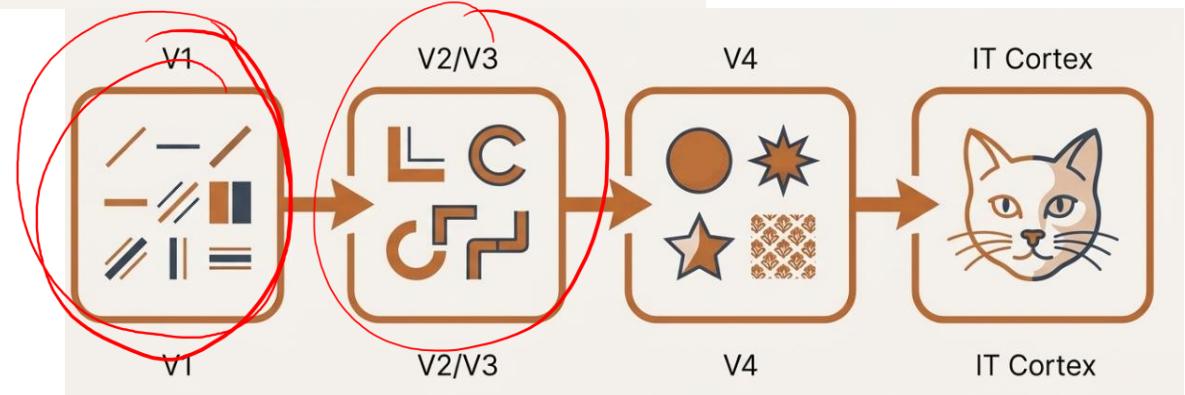


> “*No neuron is guaranteed to exist – learn features that stand alone.*”

How Brain Processes Images -> CNN



WTF. --



WTF

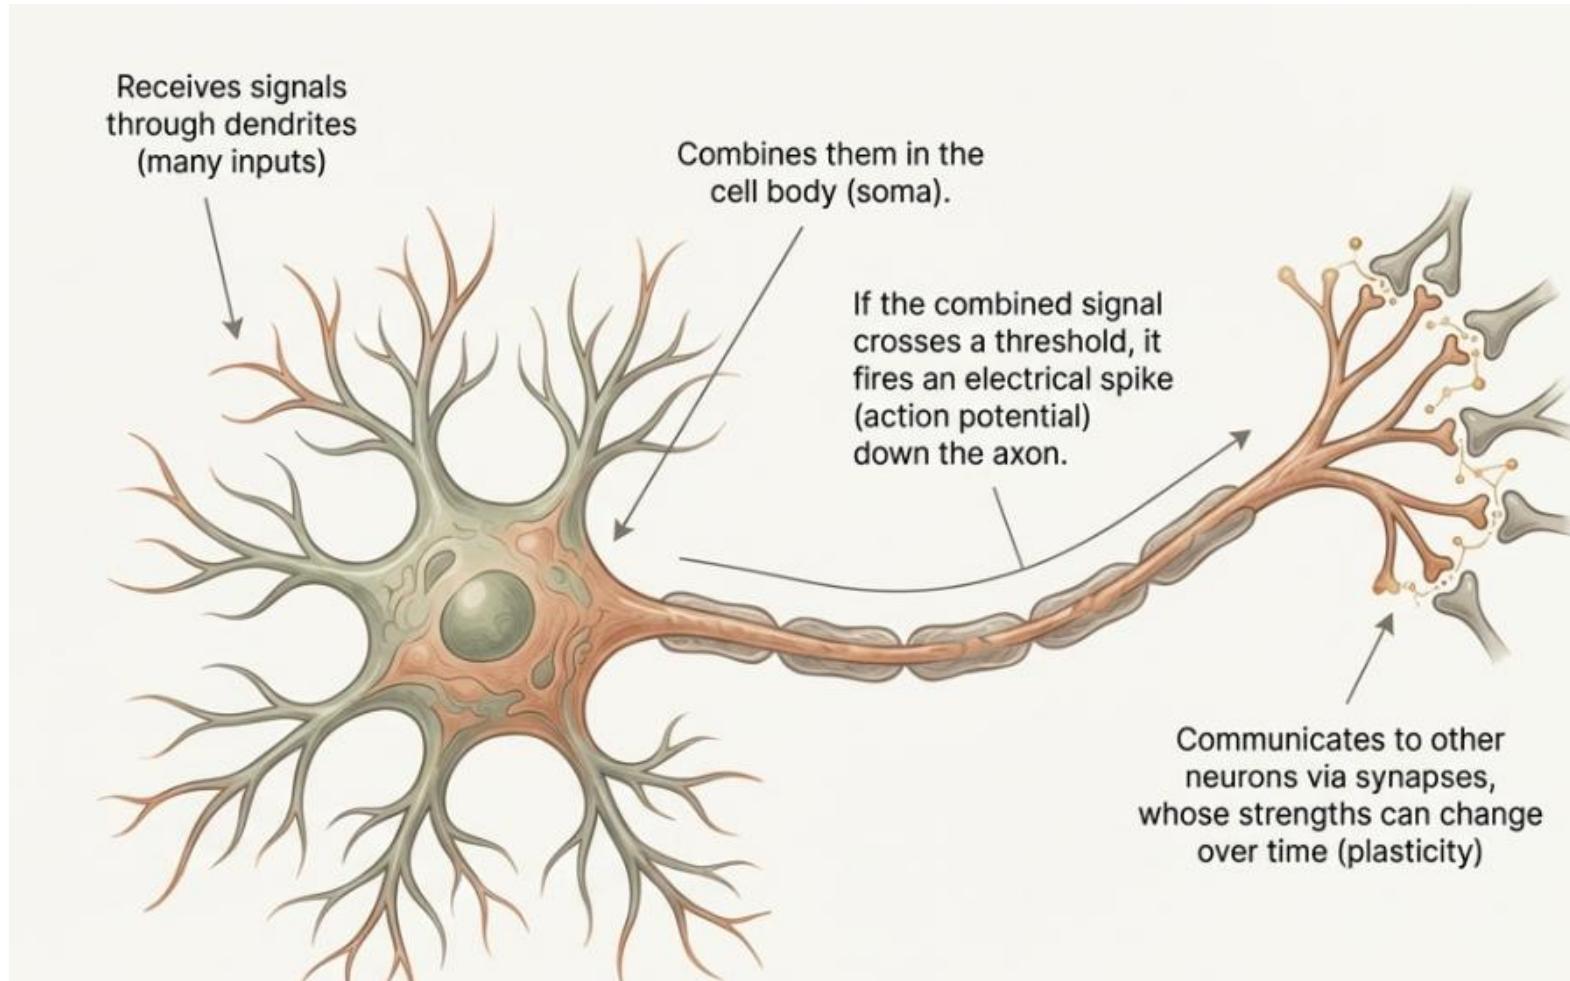
Thank You

<https://tinyurl.com/dlframeworks>

<https://github.com/sakharamg/DeepLearningFrameworks>

Appendix

Neuron in Brain



Log Softmax

0.0000302941
0.0002238446
0.0000020359
0.9954517644
0.0001500475
0.0000061163
0.0000674207
0.0003690573
0.0036810450
0.0000183743

-10.4045575839
-8.4045584969
-13.1045725764
-0.0045586103
-8.8045586473
-12.0045532194
-9.6045584655
-7.9045586414
-5.6045585997
-10.9045576088